

RKNN3 Runtime C API Reference

Document ID: RK-YH-YF-425

Version: V1.1.0

Date: 2026-08-03

Confidentiality: ☐Top Secret ☐Secret ☐Internal ☒Public

DISCLAIMER

THIS DOCUMENT IS PROVIDED "AS IS". ROCKCHIP ELECTRONICS CO., LTD. ("ROCKCHIP") DOES NOT PROVIDE ANY WARRANTY OF ANY KIND, EXPRESSED, IMPLIED OR OTHERWISE, WITH RESPECT TO THE ACCURACY, RELIABILITY, COMPLETENESS, MERCHANTABILITY, FITNESS FOR ANY PARTICULAR PURPOSE OR NON-INFRINGEMENT OF ANY REPRESENTATION, INFORMATION AND CONTENT IN THIS DOCUMENT. THIS DOCUMENT IS FOR REFERENCE ONLY. THIS DOCUMENT MAY BE UPDATED OR CHANGED WITHOUT ANY NOTICE AT ANY TIME DUE TO THE UPGRADES OF THE PRODUCT OR ANY OTHER REASONS.

Trademark Statement

"Rockchip", "瑞芯微", "瑞芯" shall be Rockchip's registered trademarks and owned by Rockchip. All the other trademarks or registered trademarks mentioned in this document shall be owned by their respective owners.

All rights reserved. ©2026. Rockchip Electronics Co., Ltd.

Beyond the scope of fair use, neither any entity nor individual shall extract, copy, or distribute this document in any form in whole or in part without the written approval of Rockchip.

Rockchip Electronics Co., Ltd.

Address: No.18, Software Park, Tongpan Road, Fuzhou, Fujian, PRC

Website: www.rock-chips.com

Customer service Tel: +86-4007-700-590

Customer service Fax: +86-591-83951833

Customer service e-Mail: fae@rock-chips.com

Preface

Overview

This document is the Rockchip RKNN3 Runtime C API Reference Manual.

Target Audience

This document (this guide) is primarily intended for the following engineers:

- Technical Support Engineers
- Software Development Engineers

Revision History

Version	Author	Date	Description	Approved By
V0.1.0	HPC Team	2025-05-12	Initial version	Vincent
V0.2.0	HPC Team	2025-08-04	Updated API interface descriptions	Vincent
V0.3.0b0	HPC Team	2025-09-10	Added description for rknn3_set_internal_mem interface; added KV cache storage methods and data type definitions; improved LLM configuration struct content	Vincent
V0.4.0b0	HPC Team	2025-11-03	Updated rknn3_vocab_info and rknn3_llm_multimodal_tensor structures and related interface descriptions	Vincent
V0.5.0	HPC Team	2025-11-29	Added post-processing related query commands and custom operator plugin mechanism; adjusted output callback function interface; added YUV input related structures and interface descriptions	Vincent
V1.0.0	HPC Team	2025-12-22	Added performance and memory analysis interfaces	Vincent
V1.0.4	HPC Team	2026-05-09	Added streaming weight loading interface; added rknn3_create_mem_from_fd interface	Vincent
V1.1.0	HPC Team	2026-08-03	Added rknn3_set_weight_mem, rknn3_set_kvcache_group, rknn3_session_pause, rknn3_session_resume, rknn3_mem_sync_range and other interfaces; restructured chapters to match the reference manual format; added multimodal inference and external memory usage flows; synchronized with rknn3_api.h header file	Vincent

Table of Contents

RKNN3 Runtime C API Reference

Table of Contents

1. Reference Manual Overview

1.1 Applicable SDK and Runtime Versions

1.2 Supported Platforms

1.3 Header File and Runtime Library

1.4 API Stability and Compatibility

2. Compilation and Linking

2.1 Linux

2.2 Android

2.3 Cross-Compilation

3. C API Common Conventions

3.1 Context, Model, and Session Lifecycle

3.2 Memory Ownership

3.3 Input and Output Data Validity

3.4 Synchronous and Asynchronous Behavior

3.5 Thread Safety and Callbacks

3.6 Return Values and Error Handling

3.7 Platform Differences

3.8 API Deprecation Rules

4. Call Flow

4.1 General Model Inference

4.2 LLM Session Inference

4.3 Multimodal Large Model Inference

4.4 External Memory Usage Flow

4.4.1 Weight Memory

4.4.2 Internal Memory

4.4.3 KV Cache Memory

5. RKNN3 C API Description

5.1. Device and Context Management

`rknn3_init`

`rknn3_destroy`

`rknn3_find_devices`

`rknn3_set_input_name_alias`

5.2. Model Loading and Initialization

`rknn3_load_model_from_path`

`rknn3_load_model_from_data`

`rknn3_model_init`

`rknn3_load_weight_chunk`

5.3. Query Interface

`rknn3_query`

5.4. Inference Execution

`rknn3_run`

`rknn3_set_shape`

5.5. Memory Management

`rknn3_create_mem_from_fd`

`rknn3_create_mem`

`rknn3_destroy_mem`

`rknn3_mem_sync`

`rknn3_mem_sync_range`

`rknn3_set_kvcache_mem`

`rknn3_set_internal_mem`

`rknn3_set_weight_mem`

`rknn3_set_kvcache_group`

5.6. Utility Functions

`rknn3_get_type_string`

`rknn3_get_qnt_type_string`

`rknn3_get_layout_string`

5.7. Model Decryption

`rknn3_set_decrypt_key_from_path`

`rknn3_set_decrypt_key_from_data`

5.8. LLM Session

`rknn3_session_init`

`rknn3_session_destroy`

`rknn3_session_run`

`rknn3_session_run_async`

`rknn3_session_stop`

`rknn3_session_pause`

`rknn3_session_resume`

`rknn3_session_query_state`

5.9. LLM Configuration and Callbacks

`rknn3_session_set_callback`

`rknn3_session_set_chat_template`

`rknn3_session_set_llm_param`

`rknn3_session_set_function_tools`

`LLMResultCallback`

`LLMSamplingCallback`

`LLMGetEmbedCallback`

`LLMTokenizerCallback`

`LLMOutputCallback`

`LLMInputCallback`

5.10. KV Cache Management

`rknn3_session_set_kvcache_policy`

`rknn3_session_clear_kvcache`

`rknn3_session_load_kvcache_from_path`

`rknn3_session_load_kvcache_from_data`

`rknn3_session_save_kvcache`

5.11. LoRA Adapters

`rknn3_lora_init`

`rknn3_lora_init_from_data`

`rknn3_lora_load`

`rknn3_lora_unload`

`rknn3_lora_enable`

`rknn3_lora_disable`

`rknn3_session_enable_lora`

`rknn3_session_disable_lora`

5.12. Profiling and Debugging

`rknn3_dump_features`

`rknn3_profile_ops`

`rknn3_profile_mem`

5.13. Custom Operators

`rknn3_register_custom_ops_plugins`

`rknn3_register_custom_op_func`

5.14. Image Processing

`rknn3_im_mem_create`

rknn3_im_mem_destroy

rknn3_im_cvt_color

5.15. Status Codes

5.16. Constants

5.16.1. Tensor Related Constants

5.16.2. Device Related Constants

5.17. Enumerations

5.17.1. Tensor and Query Enumerations

rknn3_query_cmd

rknn3_tensor_type

rknn3_tensor_qnt_type

rknn3_tensor_layout

5.17.2. Memory and Core Enumerations

rknn3_mem_alloc_flags

rknn3_mem_sync_mode

rknn3_mem_type

rknn3_core_mask

5.17.3. KV Cache Enumerations

rknn3_kvcache_policy

rknn3_kvcache_clear_policy

rknn3_kvcache_dtype

rknn3_kvcache_store_method

rknn3_attention_type

5.17.4. LLM Enumerations

rknn3_llm_input_type

LLMCallState

rknn3_llm_task_type

LLMOutputCallbackState

5.17.5. Image Processing Enumerations

rknn3_im_fmt

rknn3_im_proc_flag

5.17.6. Custom Operator Enumerations

rknn3_op_target_type

rknn3_op_plugin_type

5.18. Structures

5.18.1. Tensor Related Structures

rknn3_input_output_num
rknn3_quant_info
rknn3_tensor_attr
rknn3_tensor_mem
rknn3_tensor

5.18.2. Model and Configuration Structures

rknn3_sdk_version
rknn3_core_mem_size
rknn3_custom_string
rknn3_config
rknn3_shape_info
rknn3_shape_config
rknn3_allocation_info

5.18.3. Device Related Structures

rknn3_init_extend
rknn3_node_mem_info
rknn3_dev_mem_info
rknn3_device
rknn3_devices

5.18.4. LLM Related Structures

rknn3_llm_config
rknn3_vocab_info
rknn3_llm_extend_param
rknn3_sampling_params
rknn3_llm_param
rknn3_lora
rknn3_llm_multimodal_tensor
rknn3_llm_tensor
rknn3_llm_input
rknn3_llm_infer_param
RKLLMResult
RKLLMCallback
RKLLMRunState

5.18.5. KV Cache Related Structures

rknn3_kvcache_policy_param
rknn3_attention_kvcache_lens

[rknn3_kvcache_len_group_info](#)

5.18.6. Custom Operator Structures

[rknn3_custom_op_attr](#)

[rknn3_custom_op_context](#)

[rknn3_custom_op](#)

5.18.7. Image Processing Structures

[rknn3_im_rect](#)

[rknn3_im_metadata](#)

[rknn3_im_mem](#)

6. RKNN3 C API Change Log

[v0.3.0 -> v0.4.0 Session API Main Changes \(2025-11-03\)](#)

[v0.4.0b0 -> v0.5.0 Main Changes \(2025-11-29\)](#)

[v0.5.0 -> v1.0.0 Main Changes \(2025-12-22\)](#)

[v1.0.0 -> v1.0.4 Main Changes \(2026-05-09\)](#)

[v1.0.4 -> v1.1.0 Main Changes \(2026-08-03\)](#)

1. Reference Manual Overview

The RKNN3 C API is the C language interface for the RKNN3 Runtime. Developers use C/C++ to develop applications and deploy model inference through the RKNN3 C API. This document describes the functions, data structures, status codes, and other definitions of the RKNN3 C API.

1.1 Applicable SDK and Runtime Versions

This document corresponds to RKNN3 Runtime release version V1.1.0 (see Revision History). The interfaces described in this document are based on version V1.1.0.

1.2 Supported Platforms

This document applies to the following hardware platforms:

Coprocessor Platforms (NPU accelerators, connected to the host SoC via PCIe/USB):

- RK1820
- RK1828

General SoC Platforms (integrated NPU, applications run directly on the SoC):

- RK3572

1.3 Header File and Runtime Library

The RKNN3 SDK supports two deployment architectures:

Coprocessor mode: The application is compiled and executed on the host SoC, with the coprocessor (RK1820/RK1828) serving only as an NPU acceleration unit. The host SoC connects to the coprocessor via a high-speed PCIe/USB interface, such as an RK3588 connected to an RK1820/RK1828. In this mode, the versions of `librknn3_api`, `rknn3_transfer_proxy`, and the coprocessor firmware must be consistent.

General SoC mode: The application is compiled and executed directly on the RK3572, with the NPU being an integrated unit within the SoC, requiring no coprocessor communication. In this mode, the versions of `librknn3_api` and `rknn3_server` must be consistent.

In both modes, developers need to include the `rknn3_api.h` header file and link with the `librknn3_api.so` runtime library when compiling an RKNN3 application.

The RKNN3 C API is defined in the `rknn3_api.h` header file. It is organized by function into basic inference interfaces for general deep learning model deployment (e.g., `rknn3_init`, `rknn3_run`) and a series of session interfaces specifically optimized for large language models (e.g., `rknn3_session_init`, `rknn3_session_run`). The latter provides functionalities such as KV cache management, LoRA support, and sampling strategy configuration.

The RKNN3 SDK involves three dynamic libraries at runtime, with the following relationships:

Library Name	Description
<code>librknn3_api.so</code>	API entry layer, linked by the application. Automatically loads the corresponding backend library based on the deployment mode at runtime

Library Name	Description
librknn3_api_rkcp.so	Coprocessor mode backend library. Communicates with rknn3_transfer_proxy via localsocket, then forwards to the coprocessor via PCIe/USB for inference execution
librknn3_api_native.so	General SoC mode backend library. Directly accesses the SoC's integrated NPU without coprocessor communication

Coprocessor mode call chain:

```
APP -> librknn3_api.so -> librknn3_api_rkcp.so -> rknn3_transfer_proxy -> PCIe/USB
-> Coprocessor(RK1820/RK1828)
```

General SoC mode call chain:

```
APP -> librknn3_api.so -> librknn3_api_native.so -> NPU(integrated in SoC)
```

The application only needs to link `librknn3_api.so`; at runtime, this library automatically detects the deployment mode and loads `librknn3_api_rkcp.so` or `librknn3_api_native.so`. During deployment, ensure that the corresponding backend library is located in the same directory as `librknn3_api.so` or in the system library search path. The application can also directly link `librknn3_api_rkcp.so` or `librknn3_api_native.so` depending on the actual scenario.

There are two ways to obtain the runtime version number:

- Run the `strings` command on the corresponding runtime library and filter the version string, e.g., `strings librknn3_api.so | grep "version:"` (for coprocessor mode use `librknn3_api_rkcp.so`, for general SoC mode use `librknn3_api_native.so`) to read the runtime version from the version string;
- Query the actual version at runtime via `RKNN3_QUERY_SDK_VERSION`. The result is written to the `api_version` and `drv_version` fields of the `rknn3_sdk_version` structure (see the `rknn3_query` interface in [5. Function Reference](#)).

1.4 API Stability and Compatibility

- **Version consistency requirement:** The versions of `librknn3_api`, `rknn3_transfer_proxy` (coprocessor mode) and the coprocessor firmware, or `rknn3_server` (general SoC mode), must be consistent. If versions are incompatible, `RKNN3_ERR_INCOMPATIBLE_VERSION` is returned.
- **Platform limitations:** Some interfaces are only available on specific platforms. For example, `rknn3_set_weight_mem` only supports RK3572 in the current version.

2. Compilation and Linking

When compiling an RKNN3 application, developers need to select the appropriate cross-compilation toolchain based on the target platform's hardware architecture and system type, and ensure that the RKNN3 API library for the corresponding platform is correctly linked. The RKNN3 C API call flow is consistent across coprocessor mode and general SoC mode, with the only difference being the underlying communication method (coprocessor mode communicates with the coprocessor via PCIe/USB, while general SoC mode directly accesses the integrated NPU).

2.1 Linux

When compiling an RKNN3 application on Linux (including Linux on the host SoC), you need to include the `rknn3_api.h` header file and link with `librknn3_api.so`. The RKNN3 runtime library is a prebuilt artifact provided with the SDK; developers do not need to compile the runtime library themselves and only need to compile their own application using the corresponding cross-compilation toolchain. Compilation example:

```
aarch64-linux-gnu-gcc app.c -I/path/to/include -L/path/to/lib -lrknn3_api -o app
```

2.2 Android

When compiling an RKNN3 application on Android, use the Android NDK toolchain with the target architecture arm64-v8a. The RKNN3 runtime library is a prebuilt artifact provided with the SDK; developers do not need to compile the runtime library themselves and only need to compile their own application using the NDK toolchain. Include `rknn3_api.h` and link with `librknn3_api.so` when compiling. Compilation example:

```
aarch64-linux-android24-clang app.c -I/path/to/include -L/path/to/lib -lrknn3_api -o app
```

2.3 Cross-Compilation

In coprocessor mode, the application is compiled and executed on the host SoC, requiring the cross-compilation toolchain corresponding to the host SoC. For instructions on downloading and installing the cross-compiler for the host SoC, please refer to the [< Rockchip_RK182X_Quick_Start_RKNN3_SDK >](#) document. In general SoC mode, the application is compiled and executed directly on the RK3572, using the toolchain corresponding to RK3572.

Regardless of the mode, when compiling you must:

1. Include the directory containing `rknn3_api.h` in the header file search path;
2. Include the directory containing `librknn3_api.so` in the library search path, and link with `-lrknn3_api`;
3. Select the cross-compilation toolchain for the corresponding platform based on the deployment architecture.

3. C API Common Conventions

This chapter describes common conventions for using the RKNN3 C API, including object lifecycle, memory ownership, synchronous/asynchronous behavior, thread safety, return values and error handling, platform differences, and deprecation rules. For specific constraints of each function, see [5. Function Reference](#).

3.1 Context, Model, and Session Lifecycle

The RKNN3 API is organized around three core objects: runtime context (context), model (model), and session (session). Their typical lifecycle is as follows:

1. **Context initialization:** Call `rknn3_init` to create a context; this must be called before loading and running a model. Call `rknn3_destroy` to release when no longer needed; each context is destroyed only once.
2. **Model loading and initialization:** Load the model via `rknn3_load_model_from_path` / `rknn3_load_model_from_data`, then call `rknn3_model_init` to complete model initialization configuration.
3. **Streaming weight loading (optional):** When the weight parameter is NULL during model loading, streaming weight loading mode is entered. After `rknn3_model_init`, weight data must be uploaded in segments via `rknn3_load_weight_chunk`; when the cumulative uploaded bytes reach the total weight size, the streaming state is automatically completed.
4. **Session initialization (LLM scenarios):** Call `rknn3_session_init` to create a session; call `rknn3_session_destroy` to destroy the session when no longer needed.
5. **Resource release order:** First destroy the session (`rknn3_session_destroy`), then destroy the context (`rknn3_destroy`).

Additionally, some interfaces have requirements on calling timing:

- Model decryption key setting (`rknn3_set_decrypt_key_from_*`) must be called before `rknn3_load_model_from_*`.
- `rknn3_set_kvcache_mem` must be called after `rknn3_model_init`.
- User-allocated internal memory (`rknn3_set_internal_mem`) and weight memory (`rknn3_set_weight_mem`) require setting `config.user_mem_internal` / `config.user_mem_weight` to 1 before `rknn3_model_init`, then calling the corresponding setting interface after `rknn3_model_init`.
- LoRA initialization (`rknn3_lora_init` / `rknn3_lora_init_from_data`) must be called after `rknn3_model_init` and before the session uses LoRA.

3.2 Memory Ownership

- **API-allocated memory:** Tensor memory allocated via `rknn3_create_mem` is owned by the caller and must be released via `rknn3_destroy_mem` when no longer needed; avoid double-freeing.
- **User-allocated memory:** User-allocated memory accepted by `rknn3_set_kvcache_mem`, `rknn3_set_internal_mem`, and `rknn3_set_weight_mem` remains owned by the caller;

after use, the caller must release it via `rknn3_destroy_mem`.

- **External memory wrapping:** `rknn3_create_mem_from_fd` associates already-allocated memory (represented by a file descriptor) with the RKNN3 framework to achieve zero-copy; the wrapped `rknn3_tensor_mem` object still needs to be released via `rknn3_destroy_mem`.
- **Memory in callbacks:** Tensor data pointers passed in callback functions (e.g., custom operator attribute data, output tensors) are typically not owned by the caller and are only valid within the current callback scope; the caller should not cache or free them.

3.3 Input and Output Data Validity

- The `inputs` / `outputs` arrays of the inference interface (`rknn3_run`) and the `rknn3_tensor` and memory they point to must remain valid during the call, and the input/output tensor memory must be allocated by the user in advance.
- After the CPU writes to input memory and before the NPU accesses it, call `rknn3_mem_sync` (`RKNN3_MEMORY_SYNC_TO_DEVICE`) to synchronize CPU cache data to memory; after the NPU writes to output memory and before the CPU reads it, call `rknn3_mem_sync` (`RKNN3_MEMORY_SYNC_FROM_DEVICE`) to invalidate the CPU cache and re-read from memory. Failure to synchronize may cause abnormal inference results.
- Tensor pointers obtained in callbacks (e.g., `output_tensors` of `LLMOutputCallback`, `data` of `rknn3_custom_op_attr`) are only valid within the callback scope and should not be used after the callback returns.

3.4 Synchronous and Asynchronous Behavior

- **General model synchronous inference:** `rknn3_run` is a synchronous interface that blocks until completion during inference.
- **LLM session inference:** The inference results of `rknn3_session_run` / `rknn3_session_run_async` are obtained asynchronously via callback functions (see [5.9 LLM Configuration and Callbacks](#)).
- **Memory synchronization:** `rknn3_mem_sync` / `rknn3_mem_sync_range` are used for data consistency synchronization between the CPU and the device; `rknn3_mem_sync_range` only synchronizes the specified `[offset, offset+length)` range, suitable for partial update scenarios.

3.5 Thread Safety and Callbacks

- Concurrent use of the same `rknn3_context` or `rknn3_session` requires the caller to ensure thread safety; avoid multiple threads operating on the same handle simultaneously.
- LLM inference results are returned asynchronously via callbacks, which are triggered during inference; callbacks should not perform long-blocking operations or recursively initiate inference on the same session, so as not to affect the inference flow.
- Temporary data passed in callbacks (tensor pointers, attribute data, etc.) is only valid within the callback scope; if it needs to be retained, the caller should copy it.

3.6 Return Values and Error Handling

- Most interfaces return `int` status codes: **0 indicates success, negative values indicate errors**. A few interfaces return pointers (e.g., `rknn3_create_mem`, `rknn3_create_mem_from_fd`, `rknn3_session_init`); they return `NULL` on failure.
- For the complete status code definitions, see [5.15 Status Codes](#).
- The caller should check the return value of each call; when handling errors, use the status code to locate the problem (e.g., `RKNN3_ERR_INCOMPATIBLE_VERSION` indicates component version mismatch, `RKNN3_ERR_WEIGHT_LOAD_NOT_PREPARED` indicates streaming weight loading is not prepared, etc.).

3.7 Platform Differences

- **Deployment mode differences:** Coprocessor mode communicates with the coprocessor via PCIe/USB, requiring `librknn3_api_rkcp.so` + `rknn3_transfer_proxy` + coprocessor firmware; general SoC mode directly accesses the integrated NPU, requiring `librknn3_api_native.so` + `rknn3_server`. See [1.3 Header File and Runtime Library](#).
- **Interface availability differences:** Some interfaces are only supported on specific platforms. For example, `rknn3_set_weight_mem` only supports RK3572 in the current version.
- **Call flow consistency:** The C API call flow is consistent across coprocessor mode and general SoC mode; the only difference is the underlying communication method.

3.8 API Deprecation Rules

- For the deprecation and change history of interfaces, see [6. RKNN3 C API Change Log](#).
- All component versions must be consistent (see [1.4 API Stability and Compatibility](#)); when upgrading, `librknn3_api`, `rknn3_transfer_proxy` / `rknn3_server`, and the firmware must be updated together.

4. Call Flow

The RKNN3 C API call flow is consistent across coprocessor mode and general SoC mode, with the only difference being the underlying communication method (coprocessor mode communicates with the coprocessor via PCIe/USB, while general SoC mode directly accesses the integrated NPU). The following describes the inference flows for general model synchronous inference, general model asynchronous inference, LLM Session inference, multimodal large model inference, and external memory usage.

4.1 General Model Inference

The inference flow for non-large language models such as CNNs/ViTs using the basic API mainly includes the following steps. First, the user needs to call the `rknn3_init` interface to initialize the runtime context, and load the model using `rknn3_load_model_from_path` or `rknn3_load_model_from_data`. After the model is loaded, `rknn3_model_init` must be called for model initialization and configuration. Subsequently, the user can query the attribute information of input and output tensors through the `rknn3_query` interface and allocate memory for inputs and outputs. When performing inference, the `rknn3_run` interface is called to complete synchronous inference. Finally, `rknn3_destroy` is used to release the context resources. The inference flow for the basic API is shown in Figure 4-1.




Figure 4-1 Basic API Inference Flow

4.2 LLM Session Inference

The RKNN3 large language model inference flow typically includes the following main steps. First, the user needs to get a list of available devices by calling `rknn3_find_devices`, initialize the context with `rknn3_init`, load the model file using `rknn3_load_model_from_path`, and complete the model initialization configuration with `rknn3_model_init`. Next, the user needs to prepare resources such as the tokenizer and embeddings, and set up the `rknn3_llm_param` structure to configure session parameters. Then, call `rknn3_session_init` to initialize the session and set callback functions for the inference process via `rknn3_session_set_callback`. Before inference, a chat template can be set or the KV cache can be cleared as needed.

During the inference phase, the user constructs an `rknn3_llm_input` structure and calls the `rknn3_session_run` interface for inference. The inference results are obtained asynchronously through callback functions. After inference is complete, performance data can be collected or multi-turn dialogues can be conducted. Finally, the `rknn3_session_destroy` and `rknn3_destroy` interfaces must be called to release session and context resources, completing the entire inference flow. The large language model inference flow is shown in Figure 4-2.




Figure 4-2 Large Language Model Inference Flow

4.3 Multimodal Large Model Inference

Multimodal large model inference (supporting speech, image, text, video, and other modality inputs) is based on the LLM Session flow. The key difference from pure-text LLM inference is that, in addition to the LLM itself, independent modality encoder models (e.g., vision encoder, audio encoder) are needed to convert non-text modalities such as images and audio into embedding vectors, which are then fed into the LLM together with the text prompt to complete inference. The overall flow of multimodal large model inference is shown in Figure 4-3.




Figure 4-3 Multimodal Large Model Inference Flow

The main steps are described as follows:

1. **Model initialization:** Multimodal inference involves multiple RKNN3 contexts—one context per modality encoder (e.g., vision encoder, audio encoder) and one for the LLM. The initialization flow for each context is the same: `rknn3_init` creates the context, `rknn3_load_model_from_path` loads the model, and `rknn3_model_init` completes model initialization. For modality encoders, query input/output attributes via `RKNN3_QUERY_INPUT_ATTR / RKNN3_QUERY_OUTPUT_ATTR` and allocate memory with `rknn3_create_mem`. For the LLM, call `rknn3_session_init` to create a session, and set the chat template via `rknn3_session_set_chat_template` and callbacks via `rknn3_session_set_callback`. Additionally, auxiliary resources such as the tokenizer and embedding table must be prepared for tokenization and token embedding lookup in callbacks.
2. **Internal memory sharing:** To save memory, the contexts of modality encoders and the LLM can share the same internal memory block. During initialization, set `config.user_mem_internal` to 1 for each context. After model initialization, query the required internal memory size for each core of each context via `RKNN3_QUERY_CORE_NUMBER` and `RKNN3_QUERY_CORE_MEM_SIZE`, allocate memory with the maximum size via `rknn3_create_mem`, and set it to all contexts via `rknn3_set_internal_mem`.
3. **Modality embedding generation:** Based on the modality placeholder tags in the input prompt, run the corresponding encoders to generate embeddings. If the prompt contains `<image>`, preprocess the input image (resize, format conversion, etc.) and feed it to the vision encoder for inference (`rknn3_run`) to obtain image embeddings. If it contains `<audio>`, preprocess the input audio (e.g., extract mel features) and feed it to the audio encoder for inference to obtain audio embeddings. Video modality is similar. For modalities not present, set their count (`n_image / n_audio / n_video`) to 0 to skip.
4. **Construct multimodal input and infer:** Construct `rknn3_llm_multimodal_tensor`, set the `prompt` text, and fill each modality's embeddings and corresponding tags into the `image / audio / video` substructures (embed pointers, token counts, counts, start/end/content placeholder tags). The tags must be consistent with the model's convention, e.g., images use `<|vision_bos|> / <|vision_eos|> / <|IMAGE|>`, audio uses `<|audio_bos|> / <|audio_eos|> / <|AUDIO|>`. Use `<image> / <audio> / <video>` placeholder tags in the prompt to indicate the insertion positions of modality data; multiple modality data can be input in any order and combination. Then construct `rknn3_llm_input`, set `input_type` to `RKNN3_LLM_INPUT_MULTIMODAL`, point `multimodal_input` to the above structure, and call `rknn3_session_run` for inference. Results are obtained asynchronously via callback functions.
5. **Resource release:** After inference is complete, first release the shared internal memory, then release each modality encoder context and the LLM session and context resources.

The field definitions of the multimodal input structure `rknn3_llm_multimodal_tensor` are in [5.18.4. LLM Related Structures](#).

4.4 External Memory Usage Flow

When the user wishes to allocate and manage weight memory, internal memory, or KV cache memory themselves, the following flow can be used. The three types of external memory have

different query methods and setting timings: weight memory and internal memory require enabling the corresponding configuration flag before `rknn3_model_init`, and the size is queried via `RKNN3_QUERY_CORE_MEM_SIZE` after initialization; KV cache memory is queried via `RKNN3_QUERY_ALLOCATION_INFO` after `rknn3_model_init`. The overall usage flow for the three types of external memory is shown in Figure 4-4.




Figure 4-4 External Memory Usage Flow

Note: The query and setting of weight memory must be completed after `rknn3_model_init` and before `rknn3_run`; the same applies to internal memory. KV cache memory must be set after `rknn3_model_init`. The three types of external memory can be used individually or in combination.

Important: If the KV cache memory size (`kvcache_mem.size`) queried via `RKNN3_QUERY_ALLOCATION_INFO` is non-zero, user-allocated KV cache external memory must be set via `rknn3_set_kvcache_mem` before calling `rknn3_run` for inference; otherwise, inference will fail due to missing KV cache memory. If using the `rknn3_session_run` interface for LLM inference, there is no need to set KV cache external memory, as KV cache memory is managed internally by the session.

The specific steps for each memory type are described below:

4.4.1 Weight Memory

Weight memory is used to store model weight data. When using user-managed weight memory, the `user_mem_weight` flag must be enabled before model initialization, and the required weight memory size for each core is queried and allocated after initialization. Steps:

1. **Initialize context and load model:** Call `rknn3_init` to create a context, and load the model via `rknn3_load_model_from_path` or `rknn3_load_model_from_data`. When using user weight memory, the weight parameter of `rknn3_load_model_from_data` can be NULL (weights are subsequently uploaded via `rknn3_load_weight_chunk` in streaming mode) or the complete weight data can be passed in.
2. **Set configuration and initialize model:** Construct `rknn3_config`, set `config.user_mem_weight` to 1 and set `run_core_mask`, then call `rknn3_model_init` to complete model initialization.
3. **Query core count:** Query the number of cores via `RKNN3_QUERY_CORE_NUMBER` for subsequent per-core memory allocation.
4. **Query weight memory size:** Query the weight memory size for each core via `RKNN3_QUERY_CORE_MEM_SIZE` (the `weight_size` field of the `rknn3_core_mem_size` structure).
5. **Allocate and set weight memory:** For each core, call `rknn3_create_mem(ctx, weight_size, core_id, RKNN3_FLAG_MEMORY_CACHEABLE)` to allocate weight memory, collect into an array, and call `rknn3_set_weight_mem(ctx, mem[], n_core)` to set (currently only supports RK3572).
6. **Streaming weight loading (optional):** If the weight parameter was NULL in step 1, upload weight data in chunks via `rknn3_load_weight_chunk(ctx, offset, data, size)` until all weights are uploaded before inference.
7. **Release resources:** After inference is complete, release each weight memory via `rknn3_destroy_mem`, then call `rknn3_destroy` to release the context.

4.4.2 Internal Memory

Internal memory is used to store intermediate tensors and other data during model inference. When using user-managed internal memory, the `user_mem_internal` flag must also be enabled

before model initialization, and the required internal memory size for each core is queried and allocated after initialization. Steps:

1. **Initialize context and load model:** Call `rknn3_init` to create a context, and load the model via `rknn3_load_model_from_path` or `rknn3_load_model_from_data`.
2. **Set configuration and initialize model:** Construct `rknn3_config`, set `config.user_mem_internal` to 1 and set `run_core_mask`, then call `rknn3_model_init` to complete model initialization.
3. **Query core count:** Query the number of cores via `RKNN3_QUERY_CORE_NUMBER`.
4. **Query internal memory size:** Query the internal memory size for each core via `RKNN3_QUERY_CORE_MEM_SIZE` (the `internal_size` field of the `rknn3_core_mem_size` structure). If `internal_size` is 0, no allocation is needed for that core.
5. **Allocate and set internal memory:** For each core, call `rknn3_create_mem(ctx, internal_size, core_id, RKNN3_FLAG_MEMORY_CACHEABLE)` to allocate internal memory, collect into an array, and call `rknn3_set_internal_mem(ctx, mem[], n_core)` to set.
6. **Verify (optional):** Query memory allocation info via `RKNN3_QUERY_ALLOCATION_INFO` (the `internal_mem` field of the `rknn3_allocation_info` structure) to confirm that internal memory has been set correctly.
7. **Release resources:** After inference is complete, release each internal memory via `rknn3_destroy_mem`, then call `rknn3_destroy` to release the context.

Tip: In multimodal scenarios, multiple contexts (e.g., vision encoder, audio encoder, LLM) can share the same internal memory block to save memory. The approach is to enable `user_mem_internal` for each context and query the required memory size for each core, allocate a single shared memory block with the maximum size, and set it to all contexts via `rknn3_set_internal_mem` (see [4.3 Multimodal Large Model Inference](#)).

4.4.3 KV Cache Memory

KV cache memory is used to store Key/Value caches during large language model inference. Unlike weight memory and internal memory, KV cache memory does not require enabling a configuration flag before model initialization; instead, the required size for each core is queried via `RKNN3_QUERY_ALLOCATION_INFO` after `rknn3_model_init` and then allocated and set. Steps:

1. **Initialize context and load model:** Call `rknn3_init` to create a context, load the model via `rknn3_load_model_from_path` or `rknn3_load_model_from_data`, set `run_core_mask`, and call `rknn3_model_init` to complete model initialization.
2. **Query core count:** Query the number of cores via `RKNN3_QUERY_CORE_NUMBER`.
3. **Query KV cache memory size:** Query the memory allocation info for each core via `RKNN3_QUERY_ALLOCATION_INFO` (the `rknn3_allocation_info` structure), where the `kvcache_mem.size` field is the required KV cache memory size for that core and the `core_id` field is the physical ID of that core. If `kvcache_mem.size` is 0, no allocation is needed for that core.

4. **Allocate KV cache memory:** For cores that need allocation, call `rknn3_create_mem(ctx, kvcache_size, core_id, RKNN3_FLAG_MEMORY_CACHEABLE)` to allocate KV cache memory, and record the corresponding `core_id` for each memory.
5. **Set KV cache memory:** Call `rknn3_set_kvcache_mem(ctx, mem[], core_indices[], n_core)` to set, where `core_indices[]` is the array of physical core IDs corresponding to each memory, and `n_core` is the number of memories. Note that this interface requires a core ID array, which differs from the weight/internal memory interfaces (which take a core count).
6. **Release resources:** After inference is complete, release each KV cache memory via `rknn3_destroy_mem`, then call `rknn3_destroy` to release the context.

5. RKNN3 C API Description

API Functions

The RKNN3 C APIs are categorized by function as follows:

Category	Description
5.1. Device and Context Management	Device discovery, runtime context initialization and destruction.
5.2. Model Loading and Initialization	Model file loading, model initialization, and streaming weight loading.
5.3. Query Interface	Query model attributes, SDK version, memory allocation, etc.
5.4. Inference Execution	Synchronous/asynchronous inference execution and dynamic shape setting.
5.5. Memory Management	Tensor memory allocation, deallocation, synchronization, and user-defined memory setting.
5.6. Utility Functions	String conversion for tensor types, quantization types, and layout types.
5.7. Model Decryption	Model decryption key setting, must be called before model loading.
5.8. LLM Session	LLM session lifecycle management: creation, inference, pause/resume, destruction.
5.9. LLM Configuration and Callbacks	Parameter configuration and callback function type definitions for LLM sessions.
5.10. KV Cache Management	KV cache policy setting, clearing, loading, and saving.
5.11. LoRA Adapters	LoRA initialization, loading, enabling, and disabling.
5.12. Profiling and Debugging	Per-layer intermediate result dump, operator performance profiling, memory usage analysis.
5.13. Custom Operators	Custom operator plugin registration.
5.14. Image Processing	Image memory creation, destruction, and color space conversion.

5.1. Device and Context Management

Device discovery, runtime context initialization and destruction.

rknn3_init

This interface is used to initialize the RKNN3 runtime context and must be called before loading and running a model. context is an output parameter, and init_extend can be used to specify extended information such as device ID. After initialization, rknn3_destroy must be called to release resources when they are no longer needed. If initializing multiple times, ensure that resources are properly released to avoid memory leaks.

API	rknn3_init
Function	Initializes a new RKNN3 runtime context for running RKNN3 models on Rockchip NPU hardware.

API	rknn3_init
Parameters	rknn3_context* context: Pointer to the RKNN3 context handle to be initialized.
	rknn3_init_extend* init_extend: Pointer to device-specific initialization information.
Return	int Status code: see 5.15. Status Codes .
Note	After use, please call rknn3_destroy(context) to release the context resources.

Example:

```
rknn3_context ctx = 0;
rknn3_init_extend extend = {0};
// extend.device_id = "xxx"; // Specify for multi-device, can be set to NULL for
// single device
int ret = rknn3_init(&ctx, &extend);
if (ret < 0) { /* Error handling */ }
```

rknn3_destroy

This interface is used to destroy the RKNN3 runtime context and release related resources. After calling, the context handle becomes invalid and cannot be used for other operations. Destroying the same context multiple times leads to undefined behavior; ensure each context is destroyed only once.

API	rknn3_destroy
Function	Destroys the RKNN3 runtime context and releases resources.
Parameters	rknn3_context context: The RKNN3 context handle to destroy.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_destroy(ctx);
ctx = 0;
```

rknn3_find_devices

This interface is used to get the list of available RKNN3 devices on the current system. pdevs is a pointer to the device information structure. Ensure pdevs is valid before calling. A return value of 0 indicates no devices or not implemented.

API	rknn3_find_devices
Function	Gets the list of available RKNN3 devices.
Parameters	rknn3_devices* pdevs: Pointer to the structure that will receive the device list.
Return	int Status code: see 5.15. Status Codes .
Note	This function fills the provided rknn3_devices structure with information about all available RKNN3 devices on the system.

Example:

```
rknn3_devices devs;
int ret = rknn3_find_devices(&devs);
if (ret == 0 && devs.n_devices > 0) {
    printf("found %u devices, device[0]: id=%s\n", devs.n_devices,
    devs.devices[0].id);
}
```

rknn3_set_input_name_alias

This interface is used to set an alias for the model input name to adapt to the Session's built-in input names.

API	rknn3_set_input_name_alias
Function	Sets an alias mapping from the model input name to the Session input name.
Parameters	rknn3_context context: RKNN3 context handle obtained via rknn3_init.
	const char* src_input_name: The actual input name in the model (e.g., "input_embedding").
	const char* dst_input_name: The Session's built-in input name (e.g., "input_embeds").
Return	int Status code: see 5.15. Status Codes .
Note	When the model input name does not match the Session's built-in input name, call this interface to set the mapping. It can be called repeatedly to update the mapping. For example: the model input name is "input_embedding", while the Session's built-in name is "input_embeds"; call this interface to set the alias.

Example:

```
// Map the model input name to the session's expected input name
rknn3_set_input_name_alias(ctx, "input_embedding", "input_embeds");
```

5.2. Model Loading and Initialization

Model file loading, model initialization, and streaming weight loading.

rknn3_load_model_from_path

This interface is used to load an RKNN3 model from the specified model file path and weight file path into an initialized context. model_path and weight_path are the paths to the model and weight files respectively; weight_path can be NULL. When weight_path is NULL, streaming weight loading mode is entered, and the weight data must be uploaded in segments via the rknn3_load_weight_chunk interface after rknn3_model_init. Ensure the context is initialized and the paths are valid before loading. On failure, an error code is returned; on success, model initialization and inference can proceed.

API	rknn3_load_model_from_path
Function	Loads an RKNN3 model from file paths into the specified context.
Parameters	rknn3_context context: RKNN3 context handle.
	const char* model_path: Path to the RKNN3 model file.

API	rknn3_load_model_from_path
	const char* weight_path: Path to the RKNN3 weight file; pass NULL to enter streaming weight loading mode.
Return	int Status code: see 5.15. Status Codes .
Note	When weight_path is NULL, streaming weight loading mode is entered; rknn3_load_weight_chunk must be called after rknn3_model_init to upload weight data in segments.

Example:

```
// Normal mode: specify weight file path
rknn3_load_model_from_path(ctx, "/path/to/model.rknn", "/path/to/model.weight");

// Streaming weight loading mode: pass NULL for weight_path
rknn3_load_model_from_path(ctx, "/path/to/model.rknn", NULL);
```

rknn3_load_model_from_data

This interface is used to load an RKNN3 model from memory data, suitable for scenarios where the model and weights are already in memory. model_data and weight_data are memory pointers for the model and weights respectively; model_size and weight_size are the corresponding data sizes. When weight_data is NULL or weight_size is 0, streaming weight loading mode is entered, and the weight data must be uploaded in segments via the rknn3_load_weight_chunk interface after rknn3_model_init. Ensure the context is initialized and the data pointers and sizes are valid before calling. On failure, an error code is returned; on success, model initialization and inference can proceed.

API	rknn3_load_model_from_data
Function	Loads an RKNN3 model from memory data.
Parameters	rknn3_context context: RKNN3 context handle.
	const void* model_data: Pointer to the model data in memory.
	uint64_t model_size: Size of the model data in bytes.
	const void* weight_data: Pointer to the weight data in memory; pass NULL to enter streaming weight loading mode.
	uint64_t weight_size: Size of the weight data in bytes; pass 0 to enter streaming weight loading mode.
Return	int Status code: see 5.15. Status Codes .
Note	When weight_data is NULL or weight_size is 0, streaming weight loading mode is entered; rknn3_load_weight_chunk must be called after rknn3_model_init to upload weight data in segments.

Example:

```
// Load model and weights from memory
rknn3_load_model_from_data(ctx, model_data, model_size, weight_data, weight_size);

// Streaming mode: pass NULL for weight_data and 0 for weight_size
rknn3_load_model_from_data(ctx, model_data, model_size, NULL, 0);
```

rknn3_model_init

This interface is used to initialize the runtime configuration of a loaded model, including priority, timeout, core mask, etc. config is a pointer to the configuration parameters. Ensure the model is loaded and config is valid before calling. On failure, an error code is returned; on success, the model is ready for inference.

In streaming weight loading mode (i.e., when the weight parameter is NULL during model loading), rknn3_model_init only completes model configuration; weight data is not synchronized at this stage. After the function returns, the caller must upload weight data in segments via the rknn3_load_weight_chunk interface.

API	rknn3_model_init
Function	Initializes the RKNN3 model.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_config* config: Model configuration parameters.
Return	int Status code: see 5.15. Status Codes .
Note	In streaming weight loading mode, this interface does not transfer weight data; weights must be uploaded in segments via rknn3_load_weight_chunk after calling.

Example:

```
rknn3_config config = {0};
config.run_core_mask = RKNN3_NPU_CORE_0;
config.run_timeout = 0; // 0 means no timeout
int ret = rknn3_model_init(ctx, &config);
```

rknn3_load_weight_chunk

This interface is used to upload weight data in segments to the NPU device in streaming weight loading mode. The caller can write weight data segment by segment in any chunk size without maintaining a complete weight memory copy on the host side. The interface internally copies data to the correct addresses based on the weight memory allocation information of each NPU core.

Prerequisite: Must be called after rknn3_model_init returns successfully (i.e., streaming weight loading is ready). When the cumulative uploaded bytes reach the total weight size, the streaming state is automatically completed, and inference can proceed normally.

API	rknn3_load_weight_chunk
Function	Uploads a chunk of weight data to the NPU device in streaming weight loading mode.
Parameters	rknn3_context context: RKNN3 context handle.

API	rknn3_load_weight_chunk
	uint64_t weight_offset: Global byte offset of the weight data within the weight file.
	const void* data: Pointer to the weight data chunk.
	uint64_t size: Size of the data chunk in bytes.
Return	int Status code: see 5.15. Status Codes .
Note	Must be called after rknn3_model_init; weight_offset is the global offset relative to the beginning of the weight file; when the cumulative uploaded bytes reach the total weight size, the streaming state is automatically completed.

Example:

```
// Streaming weight loading: read weight file in a loop and upload in segments
FILE* fp = fopen("model.weight", "rb");
uint64_t offset = 0;
char buf[CHUNK_SIZE];
size_t n;
while ((n = fread(buf, 1, sizeof(buf), fp)) > 0) {
    rknn3_load_weight_chunk(ctx, offset, buf, n);
    offset += n;
}
fclose(fp);
```

5.3. Query Interface

Query model attributes, SDK version, memory allocation, etc.

rknn3_query

This interface is used to query various information about the model and runtime, such as input/output attributes, SDK version, etc. cmd specifies the query type; see [rknn3_query_cmd](#) for details. info is a pointer to the result buffer, and size is the buffer size. Ensure context and info are valid before calling, and size matches the corresponding structure size for the query type.

API	rknn3_query
Function	Queries RKNN3 information or status.
Parameters	rknn3_context context: The context of the RKNN3 model.
	rknn3_query_cmd cmd: The query command type.
	void* info: Pointer to the buffer for storing query results.
	uint64_t size: Size of the info buffer in bytes.
Return	int Status code: see 5.15. Status Codes .
Note	This function is used to query various information about the RKNN3 model and runtime, such as SDK version, device information, model information, etc. The specific information returned depends on the specified query command.

Example:

```
// Query the number of inputs and outputs
rknn3_input_output_num io_num;
rknn3_query(ctx, RKNN3_QUERY_IN_OUT_NUM, &io_num, sizeof(io_num));

// Query the attribute of input tensor 0
rknn3_tensor_attr attr;

memset(&attr, 0x00, sizeof(attr));
attr.index = 0;
rknn3_query(ctx, RKNN3_QUERY_INPUT_ATTR, &attr, sizeof(attr));
```

5.4. Inference Execution

Synchronous/asynchronous inference execution and dynamic shape setting.

rknn3_run

This interface is used to execute synchronous inference. inputs and outputs are the input and output tensor arrays respectively; n_inputs and n_outputs are the counts. Ensure context, inputs, and outputs are all properly allocated and configured before calling. The inference blocks until completion, and the return value is a status code. The memory for input and output tensors must be allocated by the user in advance.

API	rknn3_run
Function	Executes RKNN3 model inference.
Parameters	rknn3_context context: RKNN3 context handle obtained from rknn3_init.
	const rknn3_tensor inputs[]: Input tensor array containing input data.
	uint32_t n_inputs: Number of input tensors.
	rknn3_tensor outputs[]: Output tensor array for storing inference results.
	uint32_t n_outputs: Number of output tensors.
Return	int Status code: see 5.15. Status Codes .
Note	This function performs inference using the specified RKNN3 model. It receives input data through the inputs array and writes results to the outputs array. Input and output tensors must be properly allocated and configured before calling this function.

Example:


```

// Complete inference flow: query attributes -> allocate memory -> fill inputs ->
sync -> inference -> sync -> read outputs

// 1. Query the number of inputs and outputs
rknn3_input_output_num io_num;
rknn3_query(ctx, RKNN3_QUERY_IN_OUT_NUM, &io_num, sizeof(io_num));

// 2. Query each input/output tensor attribute, allocate memory by aligned_size and
core_id
rknn3_tensor inputs[io_num.n_input];
rknn3_tensor outputs[io_num.n_output];
for (uint32_t i = 0; i < io_num.n_input; i++) {
    inputs[i].attr = malloc(sizeof(rknn3_tensor_attr));
    inputs[i].attr->index = i;
    rknn3_query(ctx, RKNN3_QUERY_INPUT_ATTR, inputs[i].attr,
sizeof(rknn3_tensor_attr));
    inputs[i].mem = rknn3_create_mem(ctx, inputs[i].attr->aligned_size,
                                inputs[i].attr->core_id,
RKNN3_FLAG_MEMORY_CACHEABLE);
}
// outputs similarly ...

// 3. Fill input data, sync to device
memcpy(inputs[0].mem->virt_addr, input_data, inputs[0].attr->aligned_size);
rknn3_mem_sync(ctx, inputs[0].mem, RKNN3_MEMORY_SYNC_TO_DEVICE);

// 4. Execute inference
rknn3_run(ctx, inputs, io_num.n_input, outputs, io_num.n_output);

// 5. Sync output to CPU, read results
rknn3_mem_sync(ctx, outputs[0].mem, RKNN3_MEMORY_SYNC_FROM_DEVICE);
// Read inference results from outputs[0].mem->virt_addr ...

// 6. Free memory
for (uint32_t i = 0; i < io_num.n_input; i++) {
    rknn3_destroy_mem(ctx, inputs[i].mem);
    free(inputs[i].attr);
}
// outputs similarly ...

```

rknn3_set_shape

This interface is used to set the current shape for a dynamic input model. shape_id is the predefined shape ID in the model. Ensure the context is valid and shape_id is legitimate before calling. Only dynamic input models need to call this; static models do not require shape setting.

API	rknn3_set_shape
Function	Sets the model shape for dynamic input.
Parameters	rknn3_context context: RKNN3 model context handle.
	int32_t shape_id: The ID of the shape to set.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Dynamic shape model: set the current shape ID
rknn3_set_shape(ctx, shape_id);
```

5.5. Memory Management

Tensor memory allocation, deallocation, synchronization, and user-defined memory setting.

rknn3_create_mem_from_fd

This interface allows the user to associate already-allocated memory (represented by a file descriptor) with the RKNN3 inference framework, creating a tensor memory object from the file descriptor for efficient memory reuse and zero-copy data transfer. Ensure the context is valid and the relevant parameters (such as fd, size, etc.) are correct before calling. The allocated memory must be released via rknn3_destroy_mem when no longer needed to avoid memory leaks.

API	rknn3_create_mem_from_fd
Function	Creates a tensor memory object from a file descriptor.
Parameters	rknn3_context context: RKNN3 context handle.
	int32_t fd: File descriptor of the memory.
	void* virt_addr: Virtual address of the memory.
	uint64_t size: Size of the memory in bytes.
	uint64_t offset: Offset from the start of the memory referenced by fd.
Return	rknn3_tensor_mem* Pointer to the created tensor memory object, or NULL if creation fails.

Example:

```
// Create tensor memory from an existing fd (zero-copy scenario)
rknn3_tensor_mem* mem = rknn3_create_mem_from_fd(ctx, fd, virt_addr, size, offset);
```

rknn3_create_mem

This interface is used to allocate or register memory for a tensor. Ensure the context is valid before calling, and the core_id of input and output tensors should be set to the core_id member of the rknn3_tensor_attr structure obtained from the rknn3_query interface. The allocated memory must be released via rknn3_destroy_mem when no longer needed to avoid memory leaks.

API	rknn3_create_mem
Function	Creates a memory buffer for an RKNN3 tensor.
Parameters	rknn3_context context: RKNN3 context handle.
	uint64_t size: Size of the memory to allocate in bytes.
	int32_t core_id: Target NPU core ID for memory allocation.
	rknn3_mem_alloc_flags flags: Memory allocation flags that control allocation behavior.
Return	rknn3_tensor_mem* Pointer to the allocated tensor memory, or NULL if allocation fails.

API	rknn3_create_mem
Note	This function allocates memory that can be used for RKNN3 tensor operations. Memory is allocated on the specified core with the given flags.

Example:

```
// Allocate 1MB cacheable memory on NPU core 0
rknn3_tensor_mem* mem = rknn3_create_mem(ctx, 1024 * 1024, 0,
RKNN3_FLAG_MEMORY_CACHEABLE);
```

rknn3_destroy_mem

This interface is used to release tensor memory allocated or registered through the RKNN3 API. Ensure context and mem are valid before calling. Double-freeing the same memory leads to undefined behavior.

API	rknn3_destroy_mem
Function	Destroys memory allocated for an RKNN3 tensor.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_tensor_mem* mem: Pointer to the tensor memory structure to destroy.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_destroy_mem(ctx, mem);
mem = NULL;
```

rknn3_mem_sync

This interface is used to synchronize memory data between the CPU and the device. mode specifies the synchronization direction, commonly used before and after inference to ensure data consistency. Ensure context and mem are valid and mode is set correctly before calling. Failure to synchronize may cause abnormal inference results.

API	rknn3_mem_sync
Function	Synchronizes memory data between the CPU and the device.
Parameters	rknn3_context context: The context of the RKNN3 model.
	rknn3_tensor_mem* mem: Memory handle of the tensor.

API	rknn3_mem_sync
	rknn3_mem_sync_mode mode: The mode for flushing CPU cache and DDR data. RKNN3_MEMORY_SYNC_TO_DEVICE: Synchronizes CPU cache data to DDR. Typically used after CPU writes to memory and before the NPU accesses the same memory, to write back cache data to DDR. RKNN3_MEMORY_SYNC_FROM_DEVICE: Synchronizes DDR data to CPU cache. Typically used after the NPU writes to memory, to invalidate the cache so that the next CPU access to the same memory reads data from DDR. In FROM_DEVICE mode, data is transferred in blocks; the block size can be configured via the environment variable <code>MEM_SYNC_CHUNK_SIZE</code> (unit: bytes, default 2MiB). RKNN3_MEMORY_SYNC_BIDIRECTIONAL: Synchronizes CPU cache data to DDR and also causes the CPU to re-read data from DDR.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// After CPU writes, sync to device (before NPU reads)
rknn3_mem_sync(ctx, input_mem, RKNN3_MEMORY_SYNC_TO_DEVICE);

// After NPU writes, sync to CPU (before CPU reads)
rknn3_mem_sync(ctx, output_mem, RKNN3_MEMORY_SYNC_FROM_DEVICE);
```

rknn3_mem_sync_range

This interface is used to synchronize a specified range of memory data between the CPU and the device. Unlike `rknn3_mem_sync` which synchronizes the entire memory, this interface only synchronizes data in the `[offset, offset+length)` range, suitable for scenarios where only partial memory needs to be updated, reducing synchronization overhead. `offset` is the byte offset relative to `mem->virt_addr`, and `length` is the number of bytes to synchronize; `offset + length` must not exceed `mem->size`.

API	rknn3_mem_sync_range
Function	Synchronizes a specified range of memory data between the CPU and the device.
Parameters	rknn3_context context: The context of the RKNN3 model.
	rknn3_tensor_mem* mem: Memory handle of the tensor.
	rknn3_mem_sync_mode mode: Sync mode, same values as the mode parameter in rknn3_mem_sync .
	uint64_t offset: Starting byte offset relative to <code>mem->virt_addr</code> .
	uint64_t length: Number of bytes to synchronize (<code>offset + length</code> must not exceed <code>mem->size</code>).
Return	int Status code: 0 for success, negative value for error; see 5.15. Status Codes .
Note	Returns an error code when <code>offset + length</code> exceeds <code>mem->size</code> , <code>length</code> is 0, or <code>offset</code> is greater than or equal to <code>mem->size</code> .

Example:

```
// Sync only the [offset, offset+length) range to device
rknn3_mem_sync_range(ctx, mem, RKNN3_MEMORY_SYNC_TO_DEVICE, 1024, 4096);

// Sync only 4096 bytes from device to CPU
rknn3_mem_sync_range(ctx, mem, RKNN3_MEMORY_SYNC_FROM_DEVICE, 0, 4096);
```

rknn3_set_kvcache_mem

This interface is used to set KV cache memory for the specified NPU core. mem is the KV cache memory array, npu_core_indices is the core index array, and n_core is the number of cores. Ensure context, mem, and npu_core_indices are valid and n_core is greater than 0 before calling. mem and npu_core_indices must correspond one-to-one.

API	rknn3_set_kvcache_mem
Function	Sets KV cache memory.
Parameters	rknn3_context context: RKNN3 model context handle.
	rknn3_tensor_mem* mem[]: KV cache memory related information.
	int* npu_core_indices: NPU core indices.
	int n_core: Number of NPU cores to set kvcache for.
Return	int Status code: see 5.15. Status Codes .

Example:

```

// KV cache memory source: query each core's required size via
RKNN3_QUERY_ALLOCATION_INFO,
// allocate with rknn3_create_mem, then set with rknn3_set_kvcache_mem.
// Prerequisite: called after rknn3_model_init.

// 1. Query the number of cores
uint32_t core_num = 0;
rknn3_query(ctx, RKNN3_QUERY_CORE_NUMBER, &core_num, sizeof(core_num));

// 2. Query each core's memory allocation info (including kvcache size)
rknn3_allocation_info alloc_infos[core_num];
rknn3_query(ctx, RKNN3_QUERY_ALLOCATION_INFO, alloc_infos,
            sizeof(rknn3_allocation_info) * core_num);

// 3. Allocate memory based on each core's kvcache_mem.size
rknn3_tensor_mem* kvcache_mems[core_num];
int core_indices[core_num];
int n = 0;
for (uint32_t i = 0; i < core_num; i++) {
    uint64_t size = alloc_infos[i].kvcache_mem.size;
    if (size == 0) continue;
    kvcache_mems[n] = rknn3_create_mem(ctx, size, alloc_infos[i].core_id,
                                      RKNN3_FLAG_MEMORY_CACHEABLE);
    core_indices[n] = alloc_infos[i].core_id;
    n++;
}

// 4. Set KV cache memory
rknn3_set_kvcache_mem(ctx, kvcache_mems, core_indices, n);

// Free after use
for (int i = 0; i < n; i++)
    rknn3_destroy_mem(ctx, kvcache_mems[i]);

```

rknn3_set_internal_mem

Sets user-allocated internal tensor memory for multiple cores. Each mem's core_id field specifies the target core.

API	rknn3_set_internal_mem
Function	Sets user-allocated internal tensor memory for multiple cores.
Parameters	rknn3_context context: RKNN3 context.
	rknn3_tensor_mem* mem[]: Array of user-allocated memory objects.
	uint32_t n_core: Number of cores.
Return	int Status code: see 5.15. Status Codes .

Example:

```

// Internal memory source: query each core's required size via
RKNN3_QUERY_CORE_MEM_SIZE,
// allocate with rknn3_create_mem, then set with rknn3_set_internal_mem.
// Note: config.user_mem_internal must be set to 1 before rknn3_model_init.

// 1. Mark the use of user-allocated internal memory during model initialization
rknn3_config config = {0};
config.user_mem_internal = 1;
rknn3_model_init(ctx, &config);

// 2. Query the number of cores
uint32_t core_num = 0;
rknn3_query(ctx, RKNN3_QUERY_CORE_NUMBER, &core_num, sizeof(core_num));

// 3. Query each core's internal memory size
rknn3_core_mem_size core_mem_sizes[core_num];
rknn3_query(ctx, RKNN3_QUERY_CORE_MEM_SIZE, core_mem_sizes,
            sizeof(rknn3_core_mem_size) * core_num);

// 4. Allocate memory based on each core's internal_size
rknn3_tensor_mem* internal_mems[core_num];
for (uint32_t i = 0; i < core_num; i++) {
    if (core_mem_sizes[i].internal_size > 0) {
        internal_mems[i] = rknn3_create_mem(ctx, core_mem_sizes[i].internal_size,
                                            core_mem_sizes[i].core_id,
                                            RKNN3_FLAG_MEMORY_CACHEABLE);
    } else {
        internal_mems[i] = NULL;
    }
}

// 5. Set user-allocated internal memory
rknn3_set_internal_mem(ctx, internal_mems, core_num);

// Free after use
for (uint32_t i = 0; i < core_num; i++)
    if (internal_mems[i]) rknn3_destroy_mem(ctx, internal_mems[i]);

```

rknn3_set_weight_mem

Sets user-allocated weight memory for multiple cores. Each mem's core_id field specifies the target core. Note: The current version only supports RK3572.

API	rknn3_set_weight_mem
Function	Sets user-allocated weight tensor memory for multiple cores.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_tensor_mem* mem[]: Array of user-allocated memory objects; each mem's core_id field specifies the target core.
	uint32_t n_core: Number of cores.
Return	int Status code: see 5.15. Status Codes .
Note	The current version only supports RK3572.

Example:

```

// Weight memory source: query each core's weight size via
RKNN3_QUERY_CORE_MEM_SIZE,
// allocate with rknn3_create_mem, then set with rknn3_set_weight_mem.
// Note: config.user_mem_weight must be set to 1 before rknn3_model_init.
//      The current version only supports RK3572.

// 1. Mark the use of user-allocated weight memory during model initialization
rknn3_config config = {0};
config.user_mem_weight = 1;
rknn3_model_init(ctx, &config);

// 2. Query the number of cores and each core's weight memory size
uint32_t core_num = 0;
rknn3_query(ctx, RKNN3_QUERY_CORE_NUMBER, &core_num, sizeof(core_num));
rknn3_core_mem_size core_mem_sizes[core_num];
rknn3_query(ctx, RKNN3_QUERY_CORE_MEM_SIZE, core_mem_sizes,
            sizeof(rknn3_core_mem_size) * core_num);

// 3. Allocate memory based on each core's weight_size
rknn3_tensor_mem* weight_mems[core_num];
for (uint32_t i = 0; i < core_num; i++) {
    weight_mems[i] = rknn3_create_mem(ctx, core_mem_sizes[i].weight_size,
                                     core_mem_sizes[i].core_id,
                                     RKNN3_FLAG_MEMORY_CACHEABLE);
}

// 4. Set user-allocated weight memory
rknn3_set_weight_mem(ctx, weight_mems, core_num);

// Free after use
for (uint32_t i = 0; i < core_num; i++)
    rknn3_destroy_mem(ctx, weight_mems[i]);

```

rknn3_set_kvcache_group

Sets the active KV cache group for the current context. Each group corresponds to an entry in the `kvcache_buffer_len` list configured when converting the model with RKNN3 Toolkit. Each group corresponds to a different KV cache buffer size, i.e., a different maximum context length (max context len). Switching groups is equivalent to switching max context len.

The number of groups (`n_groups`) and the KV cache size of each group (`kvcache_sizes`) for each core can be queried via `RKNN3_QUERY_KVCACHE_LEN_GROUP_INFO` . `group_id` must be in the range `[0, n_groups)` . After calling, the internal KV cache buffer address is updated to the corresponding group's address, and the current context's execution ID is switched.

API	rknn3_set_kvcache_group
Function	Sets the active KV cache group for the current context, equivalent to switching max context len.
Parameters	rknn3_context context: RKNN3 context handle.
	int32_t group_id: The group ID to activate (must be in [0, n_groups) range).
Return	int Status code: see 5.15. Status Codes .

API	rknn3_set_kvcache_group
Note	Groups correspond one-to-one with the <code>kvcache_buffer_len</code> list set during Toolkit conversion; if using externally allocated KV cache memory, ensure the memory size meets the target group's requirements after switching groups, and reallocate and call <code>rknn3_set_kvcache_mem</code> if necessary.

Example:

```
// After querying group info, switch to group 1 (equivalent to switching max
context len)
rknn3_kvcache_len_group_info info[core_number];
rknn3_query(ctx, RKNN3_QUERY_KVCACHE_LEN_GROUP_INFO, info,
            sizeof(rknn3_kvcache_len_group_info) * core_number);
rknn3_set_kvcache_group(ctx, 1); // Switch to group_id=1
```

5.6. Utility Functions

String conversion for tensor types, quantization types, and layout types.

rknn3_get_type_string

Gets the string representation of a tensor type.

API	rknn3_get_type_string
Function	Gets the string representation of a tensor type.
Parameters	rknn3_tensor_type type: The type of the tensor.
Return	const char* The type string.

Example:

```
const char* type_str = rknn3_get_type_string(RKNN3_TENSOR_FLOAT32);
// type_str = "FP32"
```

rknn3_get_qnt_type_string

Gets the string representation of a quantization type.

API	rknn3_get_qnt_type_string
Function	Gets the string representation of a quantization type.
Parameters	rknn3_tensor_qnt_type type: The quantization type of the tensor.
Return	const char* The quantization type string.

Example:

```
const char* qnt_str =
rknn3_get_qnt_type_string(RKNN3_TENSOR_PER_CHANNEL_SYMMETRIC);
// qnt_str = "PER_CHANNEL_SYMMETRIC"
```

rknn3_get_layout_string

Gets the string representation of a layout type.

API	rknn3_get_layout_string
Function	Gets the string representation of a layout type.
Parameters	rknn3_tensor_layout layout: The layout type of the tensor.
Return	const char* The layout type string.

Example:

```
const char* layout_str = rknn3_get_layout_string(RKNN3_TENSOR_NHWC);  
// layout_str = "NHWC"
```

5.7. Model Decryption

The RKNN SDK supports encrypting RKNN models offline and decrypting them at runtime. Model decryption key setting must be called before model loading.

Model encryption is only one part of model protection and cannot completely eliminate the risk of the model being cracked. Users are advised to prevent third parties from directly accessing the model and implement overall protection at the system level.

rknn3_set_decrypt_key_from_path

This interface is used to load an RSA-encrypted user key envelope from a file path and associate it with the context. To protect the intellectual property of user models, the RKNPU3 runtime provides a model encryption and decryption mechanism. The encryption scheme uses AES-256-CTR symmetric encryption for encrypting the model file body, and RSA-OAEP asymmetric encryption for protecting the AES key. Ensure the context is initialized and key_path points to a valid key envelope file before calling.

API	rknn3_set_decrypt_key_from_path
Function	Loads an RSA-encrypted user key envelope from a file path and associates it with the context.
Parameters	rknn3_context context: RKNN3 runtime context. const char* key_path: Path to the key envelope file.
Return	int Status code: see 5.15. Status Codes .
Note	This interface must be called before rknn3_load_model_from_path / rknn3_load_model_from_data .

Example:

```
// Must be called before rknn3_load_model_from_path  
rknn3_set_decrypt_key_from_path(ctx, "/path/to/key_envelope.bin");  
rknn3_load_model_from_path(ctx, "/path/to/model.rknn", "/path/to/model.weight");
```

rknn3_set_decrypt_key_from_data

This interface is used to load an RSA-encrypted user key envelope from memory and associate it with the context. Used for scenarios where model data comes from memory transfer. Ensure `context` is initialized, `key_data` points to valid key envelope memory, and `key_size` is the byte length of the data before calling.

API	rknn3_set_decrypt_key_from_data
Function	Loads an RSA-encrypted user key envelope from memory and associates it with the context.
Parameters	rknn3_context context: RKNN3 runtime context.
	const void* key_data: Memory address of the key envelope.
	uint64_t key_size: Length of the key envelope data in bytes.
Return	int Status code: see 5.15. Status Codes .
Note	This interface must be called before <code>rknn3_load_model_from_path</code> / <code>rknn3_load_model_from_data</code> .

Example:

```
// Load key envelope from memory, must be called before loading the model
rknn3_set_decrypt_key_from_data(ctx, key_data, key_size);
```

5.8. LLM Session

LLM session lifecycle management: creation, inference, pause/resume, destruction.

rknn3_session_init

This interface is used to initialize a new RKNN3 session, supporting advanced features such as LLM. `context` is the context, `params` is the session parameters, and `n_params` is the number of parameters. Ensure `context` and `params` are valid before calling. On success, a session handle is returned; on failure, NULL is returned. The session must be released via `rknn3_session_destroy`.

API	rknn3_session_init
Function	Initializes a new RKNN3 session with the specified parameters.
Parameters	rknn3_context context: RKNN3 context pointer for the session.
	rknn3_llm_param* params: Pointer to the rknn3_llm_param structure containing session configuration parameters.
	int n_params: Number of parameters.
Return	rknn3_session* Session handle on success, NULL on initialization failure.

Example:

```
rknn3_llm_param params[1] = {0};
params[0].max_context_len = 512;
params[0].sampling_param.top_k = 40;
params[0].sampling_param.top_p = 0.9;
// ... configure vocab_info etc. ...
rknn3_session* session = rknn3_session_init(ctx, params, 1);
```

rknn3_session_destroy

This interface is used to destroy an RKNN3 session and release related resources. session is the session pointer. After calling, the session pointer becomes invalid and cannot be used for other operations. Ensure each session is destroyed only once to avoid resource leaks or double-free.

API	rknn3_session_destroy
Function	Destroys an RKNN3 session and releases related resources.
Parameters	rknn3_session* session: The RKNN3 session pointer to destroy.
Return	int Status code: see 5.15. Status Codes .
Note	After calling this function, the session pointer becomes invalid and should not be used again.

Example:

```
rknn3_session_destroy(session);
session = NULL;
```

rknn3_session_run

Runs inference with the provided inputs and parameters. Ensure session, inputs, and param are valid before calling.

API	rknn3_session_run
Function	Runs inference with the provided inputs and parameters.
Parameters	rknn3_session* session: RKNN3 session handle pointer.
	rknn3_llm_input inputs[]: Input tensor array containing input data.
	uint32_t n_inputs: Number of input tensors provided.
	rknn3_llm_infer_param* param: Pointer to inference parameter configuration.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_llm_input input = {0};
input.input_type = RKNN3_LLM_INPUT_PROMPT;
input.llm_input.prompt = "Hello, please introduce yourself.";

rknn3_llm_infer_param param = {0};
param.keep_history = 1;
param.max_new_tokens = 128;

rknn3_session_run(session, &input, 1, &param);
```

rknn3_session_run_async

Asynchronously runs inference for a large language model session. Ensure session, inputs, and param are valid before calling. Asynchronous inference requires callbacks or wait mechanisms to obtain results.

API	rknn3_session_run_async
Function	Asynchronously runs inference for a large language model session.
Parameters	rknn3_session* session: RKNN3 session handle pointer.
	rknn3_llm_input inputs[]: Input tensor array for the model.
	uint32_t n_inputs: Number of input tensors.
	rknn3_llm_infer_param* param: Pointer to inference parameter configuration.
Return	int Status code: see 5.15. Status Codes .
Note	This function performs asynchronous inference for the given LLM session. It allows non-blocking execution of the model with the provided inputs and parameters.

Example:

```
// Asynchronous inference, results obtained via callback
rknn3_session_run_async(session, &input, 1, &param);
// Results are returned asynchronously in result_callback
```

rknn3_session_stop

Terminates the current session inference. After calling, only the current inference is terminated; new inference can be initiated subsequently. Suitable for scenarios that need to interrupt long-running inference.

API	rknn3_session_stop
Function	Terminates the current session inference.
Parameters	rknn3_session* session: RKNN3 session pointer.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Interrupt current inference (new inference can be initiated afterwards)
rknn3_session_stop(session);
```

rknn3_session_pause

Pauses the execution of an RKNN3 session. After calling, the current inference is paused and can be resumed via `rknn3_session_resume`.

API	rknn3_session_pause
Function	Pauses the execution of an RKNN3 session.
Parameters	rknn3_session* session: RKNN3 session pointer to pause.
Return	int Status code: 0 for success, negative value for error; see 5.15. Status Codes .

Example:

```
rknn3_session_pause(session);
// Inference paused, can be resumed via rknn3_session_resume
```

rknn3_session_resume

Resumes the execution of a paused RKNN3 session.

API	rknn3_session_resume
Function	Resumes the execution of an RKNN3 session.
Parameters	rknn3_session* session: RKNN3 session pointer to resume.
Return	int Status code: 0 for success, negative value for error; see 5.15. Status Codes .

Example:

```
rknn3_session_resume(session);
// Resumes inference from where it was paused
```

rknn3_session_query_state

This interface is used to query the current running state of the session, including the number of generated tokens, KV policy, etc. session is the session pointer, and state is a pointer to the state structure. Ensure session and state are valid before calling. Can be used to monitor inference progress and session state.

API	rknn3_session_query_state
Function	Queries the current state of an RKNN3 session.
Parameters	rknn3_session* session: RKNN3 session pointer to query.
	RKLLMRunState* state: Pointer to store the queried running state.
Return	int Status code: see 5.15. Status Codes .

Example:

```
RKLLMRunState state;
rknn3_session_query_state(session, &state);
printf("total=%llu, prefill=%llu, decode=%llu\n",
       state.n_total_tokens, state.n_prefill_tokens, state.n_decode_tokens);
```

5.9. LLM Configuration and Callbacks

Parameter configuration and callback function type definitions for LLM sessions.

rknn3_session_set_callback

This interface is used to set callback functions for a session, supporting custom result handling, sampling, tokenization, embedding, output tensor callbacks, etc. session is the session pointer, and callback is a pointer to the callback structure. Ensure session and callback are valid before calling. Callback functions are triggered during inference to allow users to customize the inference flow.

API	rknn3_session_set_callback
Function	Sets callback functions for an RKNN3 session.
Parameters	rknn3_session* session: RKNN3 session instance pointer.
	RKLLMCallback* callback: Pointer to the RKLLMCallback structure containing callback functions.
Return	int Status code: see 5.15. Status Codes .

Example:

```
RKLLMCallback callback = {0};
callback.result_callback = my_result_callback;
callback.result_userdata = tokenizer;
callback.sampling_callback = my_sampling_callback;
callback.sampling_userdata = &embed_info;
rknn3_session_set_callback(session, &callback);
```

rknn3_session_set_chat_template

This interface is used to set the LLM chat template, including system prompt, prefix, and suffix. session is the session handle; system_prompt, prompt_prefix, and prompt_postfix are the template contents. Ensure the session is valid and the string contents meet the model's requirements before calling. Custom templates help adjust model behavior and conversation style.

API	rknn3_session_set_chat_template
Function	Sets the LLM chat template, including system prompt, prefix, and suffix.
Parameters	rknn3_session* session: RKNN3 session handle.
	const char* system_prompt: System prompt defining the language model's context or behavior.

API	rknn3_session_set_chat_template
	const char* prompt_prefix: Prefix added before user input in the conversation.
	const char* prompt_postfix: Postfix added after user input in the conversation.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_session_set_chat_template(session,
    "You are a helpful assistant.", // system_prompt
    "User: ",                       // prompt_prefix
    "\nAssistant: "                 // prompt_postfix
);
```

rknn3_session_set_llm_param

This interface is used to set LLM-related parameters (such as sampling parameters, vocabulary info, etc.) for an RKNN3 session. Ensure session and params pointers, and the number of parameters are valid.

API	rknn3_session_set_llm_param
Function	Configures LLM-related parameters (such as sampling parameters, vocabulary info, etc.), supporting updating LLM-related parameters after Session initialization.
Parameters	rknn3_session* session: RKNN3 session handle.
	rknn3_llm_param* params: Pointer to LLM parameter array (one LLM parameter per logits output).
	int n_params: Number of LLM parameters (equals the number of logits outputs).
Return	int Status code: see 5.15. Status Codes .
Note	params must be set in output order; if the model has multiple logits outputs, they must be configured accordingly.

Example:

```
// Update sampling parameters after session initialization
rknn3_llm_param params[1] = {0};
params[0].sampling_param.temperature = 0.8;
params[0].sampling_param.top_k = 50;
rknn3_session_set_llm_param(session, params, 1);
```

rknn3_session_set_function_tools

This interface is used to set function tool descriptions (function tools) for a session, enabling the model to call tools during conversations. Ensure session and tools string are valid before calling.

API	rknn3_session_set_function_tools
Function	Sets function tool descriptions for an LLM session.
Parameters	rknn3_session* session: RKNN3 session handle.

API	rknn3_session_set_function_tools
	const char* tools: Function tool description string.
Return	int Status code: see 5.15. Status Codes .

Example:

```
const char* tools = "[{\"name\":\"get_weather\",\"description\":\"Query weather\"}]";
rknn3_session_set_function_tools(session, tools);
```

LLMResultCallback

LLM result callback function type.

Function Pointer Type	LLMResultCallback
Function	Callback function for handling LLM generation results.
Parameters	void* userdata: User-defined data pointer.
	RKLLMResult* result: Pointer to the LLM result.
	LLMCallState state: LLM call state (e.g., finished, error, etc.).
Return	int Status code: see 5.15. Status Codes .

Example:

```
int my_result_callback(void* userdata, RKLLMResult* result, LLMCallState state) {
    if (state == RKLLM_RUN_NORMAL) {
        for (int i = 0; i < result->num_tokens; i++)
            printf("token[%d] = %d\n", i, result->token_ids[i]);
    } else if (state == RKLLM_RUN_FINISH) {
        printf("generation finished\n");
    }
    return 0;
}
```

LLMSamplingCallback

LLM sampling callback function type.

Function Pointer Type	LLMSamplingCallback
Function	Callback function for handling sampling logits.
Parameters	void* userdata: User-defined data pointer.
	float16* logits: Pointer to the logits array.
	char* logits_name: Name of the logits.
Return	int: Returns the selected token ID (>=0) on success, negative value for error code; see 5.15. Status Codes .

Example:

```

int my_sampling_callback(void* userdata, float16* logits, char* logits_name) {
    // Custom sampling: simple argmax
    int max_id = 0;
    for (int i = 1; i < vocab_size; i++) {
        if (fp16_to_fp32(logits[i]) > fp16_to_fp32(logits[max_id]))
            max_id = i;
    }
    return max_id; // Return the selected token ID
}

```

LLMGetEmbedCallback

Callback function type for getting LLM embedding vectors.

Function Pointer Type	LLMGetEmbedCallback
Function	Callback function for getting LLM embedding vectors.
Parameters	void* userdata: User-defined data pointer.
	int32_t* tokens: Token ID array.
	uint64_t num_tokens: Number of tokens in the tokens array.
	void* embed: Buffer pointer for storing embedding output.
	uint64_t len: Length of the embedding buffer in bytes.
Return	int Status code: see 5.15. Status Codes .

Example:

```

int my_embed_callback(void* userdata, int32_t* tokens, uint64_t num_tokens,
                     void* embed, uint64_t len) {
    // Copy embedding results to the embed buffer
    memcpy(embed, my_embed_data, len);
    return 0;
}

```

LLMTokenizerCallback

LLM tokenizer callback function type.

Function Pointer Type	LLMTokenizerCallback
Function	Callback function for handling text tokenization.
Parameters	void* userdata: User-defined data pointer.
	const char* text: Input text string pointer.
	int32_t text_len: Text length.
	int32_t* tokens: Token ID array pointer.
	int32_t n_tokens_max: Maximum number of tokens to generate.

Function Pointer Type	LLMTokenizerCallback
Return	int: Returns the number of generated tokens (≥ 0) on success, negative value for error code; see 5.15. Status Codes .

Example:

```
int my_tokenizer_callback(void* userdata, const char* text, int32_t text_len,
                        int32_t* tokens, int32_t n_tokens_max) {
    // Call custom tokenizer to convert text to token ID array
    int n = my_tokenize(text, text_len, tokens, n_tokens_max);
    return n; // Return the number of generated tokens
}
```

LLMOutputCallback

Callback function type for getting output tensors.

Function Pointer Type	LLMOutputCallback
Function	Callback function for getting output tensors.
Parameters	void* userdata: User-defined data pointer.
	rknn3_tensor* output_tensors: Output tensor array pointer.
	uint32_t n_output_tensors: Number of output tensors.
	LLMOutputCallbackState state: Output callback state.
Return	int: Returns 0 on success, non-zero on error; see 5.15. Status Codes .

Example:

```
int my_output_callback(void* userdata, rknn3_tensor* output_tensors,
                    uint32_t n_output_tensors, LLMOutputCallbackState state) {
    if (state == RKLLM_OUTPUT_CALLBACK_DECODE_FINISHED) {
        for (uint32_t i = 0; i < n_output_tensors; i++) {
            float16* data = (float16*)output_tensors[i].mem->virt_addr;
            // Process output tensor data ...
        }
    }
    return 0;
}
```

LLMInputCallback

Callback function type for providing custom input tensors, suitable for scenarios that require custom processing of inputs for each layer (e.g., gemma4's per_layer_inputs). When called, the framework passes the current inference's input tensor array and context parameters (`LLMInputCallbackParam`) to the callback, and the user fills in the corresponding tensor data in the callback.

LLMInputCallbackParam parameter structure:

Field	Type	Description
tokens	int32_t*	Pointer to the current input token ID array
num_tokens	int32_t	Number of tokens
shape_id	int32_t	Shape ID of the current input tensor
pos	int32_t	Position of the current input within the entire context
mrope_pos	int32_t	MRoPE position
mrope_start	int32_t	Starting position of MRoPE within the entire context
system_prompt_seqlen	int32_t	Sequence length of the system prompt
valid_system_prompt_seqlen	int32_t	Valid sequence length of the system prompt
max_position_embeddings	int32_t	Maximum position embeddings supported by the model
reserved	uint8_t[128]	Reserved field for future expansion
Function Pointer Type	LLMInputCallback	
Function	Callback function for providing custom input tensor data to the inference framework.	
Parameters	void* userdata: User-defined data pointer.	
	rknn3_tensor* input_tensors: Input tensor array pointer, allocated by the framework; the user fills in data in the callback.	
	uint32_t n_input_tensors: Number of input tensors.	
	LLMInputCallbackParam param: Context parameters for the current inference step, including token info, position info, etc.	
Return	int: Returns 0 on success, non-zero on error; see 5.15. Status Codes .	

Example:

```
int my_input_callback(void* userdata, rknn3_tensor* input_tensors,
                     uint32_t n_input_tensors, LLMInputCallbackParam param) {
    // Fill custom input tensors based on param.pos and other info
    for (uint32_t i = 0; i < n_input_tensors; i++) {
        // input_tensors[i].mem->virt_addr = ...;
    }
    return 0;
}
```

5.10. KV Cache Management

KV cache policy setting, clearing, loading, and saving.

rknn3_session_set_kvcache_policy

Sets the KV cache policy for an RKNN3 session. Applicable to models that require a caching mechanism, such as Transformers. Ensure the session is valid and the relevant parameters are correct before calling. KV cache operations must be used in coordination with the inference flow to avoid cache inconsistency.

API	rknn3_session_set_kvcache_policy
Function	Sets the KV cache policy for an RKNN3 session.
Parameters	rknn3_session* session: RKNN3 session handle.
	rknn3_kvcache_policy policy: The KV cache policy to set.
	rknn3_kvcache_policy_param* param: Parameters specifying the KV cache policy.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Use checkpoint policy
rknn3_kvcache_policy_param param = {0};
param.save_checkpoint.checkpoint_start_pos = 0;
param.save_checkpoint.checkpoint_interval = 128;
param.save_checkpoint.max_checkpoint_count = 10;
rknn3_session_set_kvcache_policy(session, RKNN3_KVCACHE_POLICY_SAVE_CHECKPOINT,
&param);
```

rknn3_session_clear_kvcache

Clears the KV cache of an RKNN3 session according to the specified policy.

API	rknn3_session_clear_kvcache
Function	Clears the KV cache of an RKNN3 session according to the specified policy.
Parameters	rknn3_session* session: RKNN3 session handle pointer.
	rknn3_kvcache_clear_policy clear: The policy for clearing KV cache, defining how to clear the cache.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Clear KV cache but keep system prompt
rknn3_session_clear_kvcache(session, RKNN3_KVCACHE_KEEP_SYSTEM_PROMPT);
```

rknn3_session_load_kvcache_from_path

Loads KV cache from the specified file path.

API	rknn3_session_load_kvcache_from_path
Function	Loads a KV cache file from the specified path into an RKNN3 session.
Parameters	rknn3_session* session: RKNN3 session pointer.
	const char* kvcache_path: Path to the KV cache file.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_session_load_kvcache_from_path(session, "/path/to/kvcache.bin");
```

rknn3_session_load_kvcache_from_data

Loads KV cache from memory data.

API	rknn3_session_load_kvcache_from_data
Function	Loads KV cache from memory data into an RKNN3 session.
Parameters	rknn3_session* session: RKNN3 session pointer.
	const void* kvcache_data: KV cache binary data.
	int64_t size: KV cache data length in bytes.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_session_load_kvcache_from_data(session, kvcache_data, kvcache_size);
```

rknn3_session_save_kvcache

Saves the KV cache to the specified path.

API	rknn3_session_save_kvcache
Function	Saves the KV cache to the specified path.
Parameters	rknn3_session* session: RKNN3 session pointer.
	const char* kvcache_path: Path to save the KV cache.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_session_save_kvcache(session, "/path/to/kvcache.bin");
```

5.11. LoRA Adapters

LoRA initialization, loading, enabling, and disabling.

rknn3_lora_init

Initializes context-level LoRA support from a LoRA weight file path and loads metadata. This interface must be called after `rknn3_model_init` and before the session uses LoRA. After calling, loaded LoRA information can be queried via `RKNN3_QUERY_LORA_NUM` and `RKNN3_QUERY_LORA_INFO`, and then the specified adapter can be loaded via `rknn3_lora_load`.

API	rknn3_lora_init
Function	Initializes LoRA support for an RKNN3 context from a LoRA weight file path and loads metadata.
Parameters	rknn3_context context: RKNN3 context handle.
	const char* lora_weight_path: LoRA weight file path.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Load LoRA metadata from file (must be called after model_init)
rknn3_lora_init(ctx, "/path/to/lora_weights.bin");

// Query loaded LoRA info
rknn3_lora_loras[RKNN3_MAX_LORA_NUM];
rknn3_query(ctx, RKNN3_QUERY_LORA_INFO, loras, sizeof(loras));
```

rknn3_lora_init_from_data

Initializes context-level LoRA support from LoRA weight data in memory and loads metadata. Suitable for scenarios where LoRA weights are already in memory. Usage is the same as `rknn3_lora_init`.

API	rknn3_lora_init_from_data
Function	Initializes LoRA support for an RKNN3 context from LoRA weight data in memory and loads metadata.
Parameters	rknn3_context context: RKNN3 context handle.
	const void* weight_data: Pointer to the LoRA weight data buffer.
	uint64_t weight_size: Buffer size in bytes.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_lora_init_from_data(ctx, lora_data, lora_size);
```

rknn3_lora_load

Loads the specified LoRA into the RKNN3 context. Before calling, LoRA initialization must be completed via `rknn3_lora_init` or `rknn3_lora_init_from_data`, and a valid `rknn3_lora` descriptor must be obtained via `RKNN3_QUERY_LORA_INFO`.

API	rknn3_lora_load
Function	Loads a LoRA into the RKNN3 context.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_lora* lora: Pointer to the LoRA descriptor to load.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Load the specified LoRA into the context (requires lora_init and querying lora descriptor first)
rknn3_lora_load(ctx, &loras[0]);
```

rknn3_lora_unload

Unloads the specified LoRA from the RKNN3 context, releasing its occupied resources. After unloading, the adapter is no longer available for any session.

API	rknn3_lora_unload
Function	Unloads a LoRA from the RKNN3 context.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_lora* lora: Pointer to the LoRA descriptor to unload.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_lora_unload(ctx, &loras[0]);
```

rknn3_lora_enable

Enables a loaded LoRA at the context level, making it effective for all inference under that context. If you need to independently control LoRA for a specific session, use `rknn3_session_enable_lora`.

Effect rule: This interface only updates the LoRA configuration once at the time of the call and does not repeatedly update it during subsequent inference. If a session has enabled a different LoRA via `rknn3_session_enable_lora`, that session will override this interface's configuration each time `rknn3_session_run` is called. Therefore, it is not recommended to use both interfaces to control LoRA simultaneously.

API	rknn3_lora_enable
Function	Enables a loaded LoRA at the RKNN3 context level.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_lora* lora: Pointer to the LoRA descriptor to enable.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Context-level enable (affects all inference)
rknn3_lora_enable(ctx, &loras[0]);
```

rknn3_lora_disable

Disables an enabled LoRA at the context level. After disabling, the adapter no longer takes effect for inference, but remains loaded and can be re-enabled via `rknn3_lora_enable`.

Effect rule: Same as `rknn3_lora_enable`. It is recommended not to mix context-level and session-level interfaces to control LoRA to avoid uncertain effect ordering.

API	<code>rknn3_lora_disable</code>
Function	Disables an enabled LoRA at the RKNN3 context level.
Parameters	<code>rknn3_context</code> context: RKNN3 context handle. <code>rknn3_lora* lora</code> : Pointer to the LoRA descriptor to disable.
Return	int Status code: see 5.15. Status Codes .

Example:

```
rknn3_lora_disable(ctx, &loras[0]);
```

`rknn3_session_enable_lora`

Enables LoRA for the specified session. LoRA weights are loaded at the context level via `rknn3_lora_load`, and different sessions share the same weights; this interface only controls whether the session applies the loaded LoRA and does not reload weights. Ensure session and lora are valid and lora has been loaded via `rknn3_lora_load` before calling.

Effect rule: The LoRA configuration enabled via this interface is repeatedly updated each time `rknn3_session_run` is called. Therefore, if `rknn3_lora_enable` is also used for context-level enabling of different LoRAs, `rknn3_session_run` will override the context-level configuration with the session-level configuration. It is recommended not to mix both interfaces for LoRA control to avoid uncertain effect ordering.

API	<code>rknn3_session_enable_lora</code>
Function	Enables a LoRA loaded in the context for the specified RKNN3 session. LoRA weights are managed by the context and can be shared across multiple sessions.
Parameters	<code>rknn3_session* session</code> : RKNN3 session pointer. <code>rknn3_lora* lora</code> : Pointer to the LoRA to enable.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Enable loaded LoRA (only affects the current session)
rknn3_session_enable_lora(session, &lora);
```

`rknn3_session_disable_lora`

Disables LoRA for the specified session. After disabling, the session's inference no longer applies LoRA, but the weights remain in the context and other sessions are unaffected. Can be re-enabled at any time via `rknn3_session_enable_lora`. Ensure session and lora are valid before calling.

Effect rule: Same as `rknn3_session_enable_lora`. It is recommended not to mix context-level and session-level interfaces to control LoRA to avoid uncertain effect ordering.

API	<code>rknn3_session_disable_lora</code>
Function	Disables LoRA for the specified RKNN3 session. This operation only affects the current session; LoRA weights remain in the context and do not affect other sessions.
Parameters	<code>rknn3_session* session</code> : RKNN3 session pointer.
	<code>rknn3_lora* lora</code> : Pointer to the LoRA to disable.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Disable LoRA for the current session (weights remain in the context)
rknn3_session_disable_lora(session, &lora);
```

5.12. Profiling and Debugging

Per-layer intermediate result dump, operator performance profiling, memory usage analysis.

`rknn3_dump_features`

Executes the model layer by layer and dumps each layer's output tensor in npy format to the specified directory. If no output tensor is provided, it is automatically allocated.

API	<code>rknn3_dump_features</code>
Function	Runs the model layer by layer and saves intermediate tensor data.
Parameters	<code>rknn3_context context</code> : RKNN3 context.
	<code>const rknn3_tensor inputs[]</code> : Input tensor array.
	<code>uint32_t n_inputs</code> : Number of input tensors.
	<code>rknn3_tensor outputs[]</code> : Optional output tensor array; when NULL or count is 0, internally allocated.
	<code>uint32_t n_outputs</code> : Number of output tensors; can be set to 0 to enable internal output allocation.
	<code>const char* dump_dir</code> : Directory path for saving npy files.
	<code>uint32_t core_mask</code> : NPU core mask to dump; when set to 0, uses the <code>run_core_mask</code> configured in <code>rknn3_model_init</code> .
	<code>const char dump_selector</code> : Dump mode string. NULL, empty string, "", or "all" means dump all operators on the selected cores.
Return	int Status code: see 5.15. Status Codes .
Note	The dump directory must be writable; only cores with bits set in <code>core_mask</code> perform feature dump, and the selected core mask must be a subset of the model's running core mask.

Example:

```
// Dump intermediate outputs of all operators to the specified directory (output
tensors auto-allocated)
rknn3_dump_features(ctx, inputs, n_inputs, NULL, 0, "/tmp/dump", 0, "all");
```

rknn3_profile_ops

Executes the model layer by layer and prints operator performance information, including time, cycles, and bandwidth. If no output tensor is provided, it is automatically allocated.

API	rknn3_profile_ops
Function	Runs the model layer by layer and prints operator performance information.
Parameters	rknn3_context context: RKNN3 context handle.
	const rknn3_tensor inputs[]: Input tensor array.
	uint32_t n_inputs: Number of input tensors.
	rknn3_tensor outputs[]: Output tensor array, can be NULL.
	uint32_t n_outputs: Number of output tensors; can be 0 to enable internal output allocation.
	uint32_t log_level: Log level (0: summary only; 1: operator details + summary; 2: operator details (with time/cycles/bandwidth) + summary).
Return	int Status code: see 5.15. Status Codes .
Note	The cores used for per-layer profiling come from the run_core_mask set in <code>rknn3_model_init</code> .

Example:

```
// Print operator-level performance info (log_level=2 includes
time/cycles/bandwidth)
rknn3_profile_ops(ctx, inputs, n_inputs, NULL, 0, 2);
```

rknn3_profile_mem

This interface is used to print memory usage information for each NPU core; must be called after `rknn3_model_init`.

API	rknn3_profile_mem
Function	Prints memory usage information for each NPU core.
Parameters	rknn3_context context: RKNN3 context handle.
Return	int Status code: see 5.15. Status Codes .
Note	Recommended to call after <code>rknn3_model_init</code> to ensure memory has been allocated.

Example:

```
// Call after model_init to print each core's memory usage
rknn3_profile_mem(ctx);
```

5.13. Custom Operators

Custom operator plugin registration.

rknn3_register_custom_ops_plugins

Registers custom operator plugins (currently supports post-processing plugins). This interface loads the plugin and registers it to the specified context.

API	rknn3_register_custom_ops_plugins
Function	Registers custom operator plugins.
Parameters	rknn3_context ctx: RKNN3 context handle.
	const char* plugin_path: Path to the plugin shared library.
	int64_t size: Reserved field, can be set to 0.
Return	int Status code: see 5.15. Status Codes .

Example:

```
// Register post-processing plugin shared library
rknn3_register_custom_ops_plugins(ctx, "/path/to/postprocess_plugin.so", 0);
```

rknn3_register_custom_op_func

Custom operator registration function pointer type, used to register custom operators within the plugin library.

Function Pointer Type	rknn3_register_custom_op_func
Function	Returns a custom operator descriptor structure pointer based on the operator index.
Parameters	int op_index: Operator index (starting from 0).
Return	rknn3_custom_op*: Pointer to the corresponding custom operator descriptor; returns NULL if the index does not exist.

Example:

```
// Exported function within the plugin library, returns custom operator descriptor
// by index
rknn3_custom_op* my_register_func(int op_index) {
    if (op_index == 0) return &my_postprocess_op;
    return NULL;
}
```

5.14. Image Processing

Image memory creation, destruction, and color space conversion.

rknn3_im_mem_create

This interface is used to create an RKNN3 image memory object `rknn3_im_mem`, typically used for image pre-processing/post-processing flows before and after model inference.

API	rknn3_im_mem_create
Function	Allocates an RKNN3 image memory object on the device.
Parameters	rknn3_context context: RKNN3 context handle.
	int32_t width: Image width in pixels.
	int32_t height: Image height in pixels.
	rknn3_im_fmt format: Image format; see rknn3_im_fmt .
	int32_t size: Image memory size in bytes; when 0, automatically calculated from width, height, and format.
	int32_t core_id: Target NPU core ID for allocation.
	rknn3_mem_alloc_flags flags: Memory allocation flags.
	rknn3_im_mem* im_mem: Output parameter, receives the created <code>rknn3_im_mem</code> object pointer.
Return	int Status code: see 5.15. Status Codes .
Note	After successful creation, the object must be released with <code>rknn3_im_mem_destroy</code> .

Example:

```
rknn3_im_mem im_mem;
// Allocate 1920x1080 RGB888 image memory
rknn3_im_mem_create(ctx, 1920, 1080, RKNN3_IM_FMT_RGB888, 0, 0,
                    RKNN3_FLAG_MEMORY_CACHEABLE, &im_mem);
```

rknn3_im_mem_destroy

This interface is used to destroy an RKNN3 image memory object created via `rknn3_im_mem_create` and release its occupied device memory.

API	rknn3_im_mem_destroy
Function	Destroys an RKNN3 image memory object and releases device memory.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_im_mem* im_mem: Pointer to the <code>rknn3_im_mem</code> object to destroy.
Return	int Status code: see 5.15. Status Codes .
Note	After calling, <code>im_mem</code> must not be used again.

Example:

```
rknn3_im_mem_destroy(ctx, &im_mem);
```

rknn3_im_cvt_color

This interface is used to perform image color format conversion on RKNN3 image memory, converting an input image from one RGB/YUV format to another before model inference.

API	rknn3_im_cvt_color
Function	Performs color space conversion between RKNN3 image memory objects.
Parameters	rknn3_context context: RKNN3 context handle.
	rknn3_im_mem* src: Source image memory object.
	rknn3_im_mem* dst: Destination image memory object.
Return	int Status code: see 5.15. Status Codes .
Note	Both <code>src</code> and <code>dst</code> must be created by <code>rknn3_im_mem_create</code> .

Example:

```
// Convert RGB888 to BGR888 (both src and dst must be created with im_mem_create first)
rknn3_im_cvt_color(ctx, &src_im_mem, &dst_im_mem);
```

5.15. Status Codes

Status Code	Value	Description
RKNN3_SUCCESS	0	Success
RKNN3_ERR_FAIL	-1	Execution failure
RKNN3_ERR_ARGUMENT_INVALID	-2	Invalid parameter
RKNN3_ERR_MODEL_INVALID	-3	Invalid model
RKNN3_ERR_CTX_INVALID	-4	Invalid context
RKNN3_ERR_RUN_TASK_FAILED	-5	Task execution failure
RKNN3_ERR_OUT_OF_MEMORY	-6	Out of memory
RKNN3_ERR_TIMEOUT	-7	Execution timeout
RKNN3_ERR_INPUT_INVALID	-8	Invalid input
RKNN3_ERR_OUTPUT_INVALID	-9	Invalid output
RKNN3_ERR_DEVICE_UNAVAILABLE	-10	Device unavailable
RKNN3_ERR_DEVICE_UNMATCH	-11	Device mismatch
RKNN3_ERR_TARGET_PLATFORM_UNMATCH	-12	Target platform mismatch
RKNN3_ERR_COMMUNICATION	-13	Communication error
RKNN3_ERR_MEM_SYNC_FAILED	-14	Memory synchronization failure
RKNN3_ERR_KEY_INVALID	-15	Invalid decryption key or RSA decryption failure
RKNN3_ERR_DECRYPT_FAILED	-16	Model decryption failed (key mismatch or data corruption)
RKNN3_ERR_WEIGHT_LOAD_NOT_PREPARED	-17	Streaming weight loading not prepared; please call <code>rknn3_model_init</code> first

Status Code	Value	Description
RKNN3_ERR_INCOMPATIBLE_VERSION	-18	Component version incompatible: librknn3_api, rknn3_transfer_proxy, and coprocessor firmware (or rknn3_server for non-coprocessor mode) must have consistent versions
RKNN3_WARN_NPU_CORE_UNUSED	-100	NPU core unused; warning only, does not affect model execution

5.16. Constants

5.16.1. Tensor Related Constants

Constant Name	Value	Description
RKNN3_MAX_DIMS	16	Maximum number of tensor dimensions
RKNN3_MAX_STRIDE_DIMS	17	Maximum number of tensor stride dimensions
RKNN3_MAX_NAME_LEN	256	Maximum length of a tensor name
RKNN3_MAX_DYNAMIC_SHAPE_NUM	512	Maximum number of dynamic shapes per input
RKNN3_MAX_LORA_NUM	32	Maximum number of LoRA adapters
RKNN3_MAX_SPECIAL_BOS_ID_NUM	64	Maximum number of special beginning-of-sequence (BOS) token IDs
RKNN3_MAX_SPECIAL_EOS_ID_NUM	64	Maximum number of special end-of-sequence (EOS) token IDs
RKNN3_MAX_KVCACHE_LEN_GROUPS	16	Maximum number of KV cache length groups
RKNN3_MAX_ATTENTION_TYPE_NUM	3	Maximum number of attention types

5.16.2. Device Related Constants

Constant Name	Value	Description
RKNN3_MAX_DEVS	64	Maximum number of devices
RKNN3_MAX_DEV_LEN	64	Maximum length of device ID/type
RKNN3_MAX_NPU_NODE_NUM	128	Maximum number of NPU nodes

5.17. Enumerations

The RKNN3 enumeration types are categorized by function as follows:

Category	Description
5.17.1. Tensor and Query Enumerations	Contains 4 enumerations
5.17.2. Memory and Core Enumerations	Contains 4 enumerations
5.17.3. KV Cache Enumerations	Contains 5 enumerations
5.17.4. LLM Enumerations	Contains 4 enumerations
5.17.5. Image Processing Enumerations	Contains 2 enumerations

Category	Description
5.17.6. Custom Operator Enumerations	Contains 2 enumerations

5.17.1. Tensor and Query Enumerations

RKNN3_QUERY_CMD

Query command enumeration, used for the `rknn3_query` function.

Enum Value	Description
RKNN3_QUERY_IN_OUT_NUM	Query the number of input and output tensors
RKNN3_QUERY_INPUT_ATTR	Query the attributes of input tensors
RKNN3_QUERY_OUTPUT_ATTR	Query the attributes of output tensors
RKNN3_QUERY_SDK_VERSION	Query SDK and driver version
RKNN3_QUERY_CORE_MEM_SIZE	Query weight and internal memory size for each core
RKNN3_QUERY_CUSTOM_STRING	Query custom string
RKNN3_QUERY_ORIGIN_INPUT_ATTR	Query the attributes of original input tensors
RKNN3_QUERY_ORIGIN_OUTPUT_ATTR	Query the attributes of original output tensors
RKNN3_QUERY_DEVICE_MEM_INFO	Query RKNN3 memory information attributes
RKNN3_QUERY_CORE_NUMBER	Query the number of cores
RKNN3_QUERY_ALLOCATION_INFO	Query allocation information
RKNN3_QUERY_DYNAMIC_SHAPE_CONFIG	Query the full dynamic shape configuration
RKNN3_QUERY_DYNAMIC_SHAPE_INFO	Query all supported shape combinations
RKNN3_QUERY_LLM_CONFIG	Query LLM-specific configuration, such as chat template
RKNN3_QUERY_POSTPROCESS_IN_OUT_NUM	Query the number of post-processing inputs and outputs; only valid when post-processing is enabled
RKNN3_QUERY_POSTPROCESS_OUTPUT_ATTR	Query post-processing output tensor attributes; only valid when post-processing is enabled
RKNN3_QUERY_POSTPROCESS_DYNAMIC_SHAPE_INFO	Query post-processing dynamic shape information; only valid when post-processing is enabled
RKNN3_QUERY_LORA_NUM	Query the number of initialized LoRAs; must be called after <code>rknn3_lora_init</code> / <code>rknn3_lora_init_from_data</code> ; result is <code>uint32_t</code>
RKNN3_QUERY_LORA_INFO	Query LoRA information list; result is written to a <code>rknn3_lora</code> array (size <code>RKNN3_MAX_LORA_NUM</code>)
RKNN3_QUERY_KVCACHE_LEN_GROUP_INFO	Query KV cache length group information (number of groups, KV cache size per group per core)

Enum Value	Description
RKNN3_QUERY_CMD_MAX	Maximum value of the query command enumeration

RKNN3_TENSOR_TYPE

Tensor data type enumeration.

Enum Value	Description
RKNN3_TENSOR_FLOAT32	float32 type
RKNN3_TENSOR_FLOAT16	float16 type
RKNN3_TENSOR_INT8	int8 type
RKNN3_TENSOR_UINT8	uint8 type
RKNN3_TENSOR_INT16	int16 type
RKNN3_TENSOR_UINT16	uint16 type
RKNN3_TENSOR_INT32	int32 type
RKNN3_TENSOR_UINT32	uint32 type
RKNN3_TENSOR_INT64	int64 type
RKNN3_TENSOR_UINT64	uint64 type
RKNN3_TENSOR_BOOL	bool type
RKNN3_TENSOR_INT4	int4 type
RKNN3_TENSOR_FLOAT8E4M3FN	float8e4m3fn type
RKNN3_TENSOR_BFLOAT16	bfloat16 type
RKNN3_TENSOR_FLOAT8E8M0	float8e8m0 type
RKNN3_TENSOR_FLOAT4E2M1	float4e2m1 (fp4) type
RKNN3_TENSOR_TYPE_MAX	Maximum value of the tensor data type enumeration

RKNN3_TENSOR_QNT_TYPE

Quantization type enumeration.

Enum Value	Description
RKNN3_TENSOR_QNT_NONE	No quantization
RKNN3_TENSOR_PER_LAYER_SYMMETRIC	Per-layer symmetric quantization
RKNN3_TENSOR_PER_LAYER_ASYMMETRIC	Per-layer asymmetric quantization
RKNN3_TENSOR_PER_CHANNEL_SYMMETRIC	Per-channel symmetric quantization
RKNN3_TENSOR_PER_CHANNEL_ASYMMETRIC	Per-channel asymmetric quantization
RKNN3_TENSOR_PER_GROUP_SYMMETRIC	Per-group symmetric quantization
RKNN3_TENSOR_PER_GROUP_ASYMMETRIC	Per-group asymmetric quantization
RKNN3_TENSOR_QNT_MAX	Maximum value of the quantization type enumeration

RKNN3_TENSOR_LAYOUT

Tensor data layout enumeration.

Enum Value	Description
RKNN3_TENSOR_UNDEFINED	Undefined
RKNN3_TENSOR_NCHW	Data layout is NCHW
RKNN3_TENSOR_NHWC	Data layout is NHWC
RKNN3_TENSOR_NC1HWC2	Data layout is NC1HWC2
RKNN3_TENSOR_CHWN	Reserved value, not currently used
RKNN3_TENSOR_HWIO	Reserved value, not currently used
RKNN3_TENSOR_OIHW	Reserved value, not currently used
RKNN3_TENSOR_OI1HWI2O2	Reserved value, not currently used
RKNN3_TENSOR_LAYOUT_MAX	Maximum value of the tensor data layout enumeration

5.17.2. Memory and Core Enumerations

RKNN3_MEM_ALLOC_FLAGS

Memory allocation flags for creating RKNN3 tensor memory.

Enum Value	Description
RKNN3_FLAG_MEMORY_FLAGS_DEFAULT	Same as RKNN3_FLAG_MEMORY_CACHEABLE
RKNN3_FLAG_MEMORY_CACHEABLE	Create cacheable memory
RKNN3_FLAG_MEMORY_NON_CACHEABLE	Create non-cacheable memory

RKNN3_MEM_SYNC_MODE

Memory synchronization mode for the rknn3_mem_sync function.

Enum Value	Description
RKNN3_MEMORY_SYNC_TO_DEVICE	Synchronizes CPU cache data to DDR. Typically used after CPU writes to memory and before the NPU accesses the same memory, to write back cache data to DDR.
RKNN3_MEMORY_SYNC_FROM_DEVICE	Synchronizes DDR data to CPU cache. Typically used after the NPU writes to memory, to invalidate the cache so that the next CPU access to the same memory reads data from DDR.
RKNN3_MEMORY_SYNC_BIDIRECTIONAL	Synchronizes CPU cache data to DDR and also causes the CPU to re-read data from DDR.

RKNN3_MEM_TYPE

Memory type enumeration.

Enum Value	Description
RKNN3_MEMORY_TYPE_NPU_DRAM	NPU DRAM memory
RKNN3_MEMORY_TYPE_EXT_DDR	External DDR memory

RKNN3_CORE_MASK

Mode for running on target NPU cores.

Enum Value	Description
RKNN3_NPU_CORE_AUTO	Default; indicates automatic scheduling, automatically running on the currently idle NPU core
RKNN3_NPU_CORE_0	Run on NPU core 0
RKNN3_NPU_CORE_1	Run on NPU core 1
RKNN3_NPU_CORE_2	Run on NPU core 2
RKNN3_NPU_CORE_3	Run on NPU core 3
RKNN3_NPU_CORE_4	Run on NPU core 4
RKNN3_NPU_CORE_5	Run on NPU core 5
RKNN3_NPU_CORE_6	Run on NPU core 6
RKNN3_NPU_CORE_7	Run on NPU core 7
RKNN3_NPU_CORE_ALL	Indicates running on all NPU cores

5.17.3. KV Cache Enumerations

RKNN3_KVCACHE_POLICY

KV cache policy.

Enum Value	Description
RKNN3_KVCACHE_POLICY_DEFAULT	Default policy, equivalent to RKNN3_KVCACHE_POLICY_RECURRENT
RKNN3_KVCACHE_POLICY_RECURRENT	Use recurrent cache policy
RKNN3_KVCACHE_POLICY_NORMAL	Use normal cache policy; KV cache is used only within max_context_len
RKNN3_KVCACHE_POLICY_SAVE_CHECKPOINT	Use checkpoint cache policy

RKNN3_KVCACHE_CLEAR_POLICY

Policy for clearing KV cache.

Enum Value	Description
RKNN3_KVCACHE_CLEAR_ALL	Clear all KV cache entries completely
RKNN3_KVCACHE_KEEP_SYSTEM_PROMPT	Clear KV cache but keep the system prompt

RKNN3_KVCACHE_DTYPE

KV cache data type. These types are internally defined data types that affect KV cache size and management; users do not need to parse them.

Enum Value	Description
RKNN3_KVCACHE_DTYPE_UNDEFINED	Undefined
RKNN3_KVCACHE_DTYPE_INT4_TO_F16	INT4 to FLOAT16
RKNN3_KVCACHE_DTYPE_INT4_TO_F8	INT4 to FLOAT8
RKNN3_KVCACHE_DTYPE_INT8_TO_F16	INT8 to FLOAT16
RKNN3_KVCACHE_DTYPE_FLOAT4_TO_F16	FLOAT4 to FLOAT16
RKNN3_KVCACHE_DTYPE_FLOAT4_TO_F8	FLOAT4 to FLOAT8
RKNN3_KVCACHE_DTYPE_FLOAT8_TO_F16	FLOAT8 to FLOAT16
RKNN3_KVCACHE_DTYPE_FLOAT8_TO_F8	FLOAT8 to FLOAT8
RKNN3_KVCACHE_DTYPE_FLOAT16	FLOAT16

RKNN3_KVCACHE_STORE_METHOD

KV cache storage method. These types are internally defined data types that affect KV cache size and management; users do not need to parse them.

Enum Value	Description
RKNN3_KVCACHE_STORE_METHOD_UNDEFINED	Undefined
RKNN3_KVCACHE_STORE_METHOD_NORMAL	Normal storage
RKNN3_KVCACHE_STORE_METHOD_GROUP_QUANT	Group quantization storage

RKNN3_ATTENTION_TYPE

LLM attention type enumeration, used for attention configuration.

Enum Value	Description
RKNN3_ATTENTION_TYPE_FULL_ATTENTION	Full attention
RKNN3_ATTENTION_TYPE_SLIDING_ATTENTION	Sliding attention
RKNN3_ATTENTION_TYPE_LINEAR_ATTENTION	Linear attention

5.17.4. LLM Enumerations

RKNN3_LLM_INPUT_TYPE

Defines the input types that can be provided to the LLM.

Enum Value	Description
RKNN3_LLM_INPUT_PROMPT	Input is a text prompt
RKNN3_LLM_INPUT_TOKEN	Input is a sequence of tokens

Enum Value	Description
RKNN3_LLM_INPUT_EMBED	Input is embedding vectors
RKNN3_LLM_INPUT_MULTIMODAL	Multimodal input
RKNN3_LLM_INPUT_AUX	Other type of input
RKNN3_LLM_INPUT_MAX	Maximum value of the input type enumeration

LLMCALLSTATE

Describes the possible states of an LLM call.

Enum Value	Description
RKLLM_RUN_NORMAL	LLM call is in normal running state
RKLLM_RUN_WAITING	LLM call is waiting for a complete UTF-8 encoded character
RKLLM_RUN_FINISH	LLM call has finished execution
RKLLM_RUN_STOP	User stopped the LLM call
RKLLM_RUN_MAX_NEW_TOKEN_REACHED	Maximum new token count has been reached
RKLLM_RUN_ERROR	An error occurred during the LLM call
RKLLM_RUN_PAUSE	LLM call has been paused
RKLLM_RUN_RESUME	LLM call has resumed from pause

RKNN3_LLM_TASK_TYPE

LLM task type enumeration.

Enum Value	Description
RKNN3_LLM_TASK_GENERATE	Generation task
RKNN3_LLM_TASK_EMBEDDING	Embedding task
RKNN3_LLM_TASK_RERANKER	Reranking task

LLMOUTPUTCALLBACKSTATE

Output callback stage enumeration.

Enum Value	Description
RKLLM_OUTPUT_CALLBACK_PREFILL_PROCESSING	Prefill stage callback processing
RKLLM_OUTPUT_CALLBACK_PREFILL_FINISHED	Prefill stage callback finished
RKLLM_OUTPUT_CALLBACK_DECODE_PROCESSING	Decode stage callback processing
RKLLM_OUTPUT_CALLBACK_DECODE_FINISHED	Decode stage callback finished

5.17.5. Image Processing Enumerations

RKNN3_IM_FMT

Image format enumeration, used for RKNN3 image-related interfaces.

Enum Value	Description
RKNN3_IM_FMT_RGB888	RGB888 format, commonly used for general image preprocessing
RKNN3_IM_FMT_BGR888	BGR888 format, commonly used for OpenCV image processing
RKNN3_IM_FMT_GRAY8	8-bit grayscale image format
RKNN3_IM_FMT_YCbCr_420_SP	YCbCr 4:2:0 semi-planar format, commonly used for video encoding/decoding
RKNN3_IM_FMT_YCrCb_420_SP	YCrCb 4:2:0 semi-planar format, commonly used for video encoding/decoding
RKNN3_IM_FMT_YCbCr_422_SP	YCbCr 4:2:2 semi-planar format, commonly used for video encoding/decoding
RKNN3_IM_FMT_YCrCb_422_SP	YCrCb 4:2:2 semi-planar format, commonly used for video encoding/decoding
RKNN3_IM_FMT_UNKNOWN	Unknown or unsupported image format

RKNN3_IM_PROC_FLAG

Image processing flag bit enumeration.

Enum Value	Description
RKNN3_IM_PROC_FLAG_NONE	No processing
RKNN3_IM_PROC_FLAG_CROP	Crop
RKNN3_IM_PROC_FLAG_RESIZE	Resize
RKNN3_IM_PROC_FLAG_FILL	Fill
RKNN3_IM_PROC_FLAG_COLOR_SPACE_CONVERT	Color space conversion
RKNN3_IM_PROC_FLAG_DECODE	Decode
RKNN3_IM_PROC_FLAG_ENCODE	Encode

5.17.6. Custom Operator Enumerations

RKNN3_OP_TARGET_TYPE

Custom operator backend target enumeration.

Enum Value	Description
RKNN3_OP_TARGET_TYPE_CPU	Run on CPU
RKNN3_OP_TARGET_TYPE_MAX	Maximum value of the enumeration

RKNN3_OP_PLUGIN_TYPE

Custom operator plugin type enumeration.

Enum Value	Description
RKNN3_OP_PLUGIN_TYPE_POSTPROCESS	Post-processing plugin
RKNN3_OP_PLUGIN_TYPE_CUSTOM_OP	Custom operator plugin (currently not supported)
RKNN3_OP_PLUGIN_TYPE_MAX	Maximum value of the enumeration

5.18. Structures

The RKNN3 structure types are categorized by function as follows:

Category	Description
5.18.1. Tensor Related Structures	Contains 5 structures
5.18.2. Model and Configuration Structures	Contains 7 structures
5.18.3. Device Related Structures	Contains 5 structures
5.18.4. LLM Related Structures	Contains 13 structures
5.18.5. KV Cache Related Structures	Contains 3 structures
5.18.6. Custom Operator Structures	Contains 3 structures
5.18.7. Image Processing Structures	Contains 3 structures

5.18.1. Tensor Related Structures

RKNN3_INPUT_OUTPUT_NUM

Information structure for RKNN3_QUERY_IN_OUT_NUM.

Member	Data Type	Description
n_input	uint32_t	Number of inputs
n_output	uint32_t	Number of outputs

RKNN3_QUANT_INFO

Quantization information structure.

Member	Data Type	Description
scale	float	Scale data
zero_point	int32_t	Zero point data

RKNN3_TENSOR_ATTR

Structure for storing RKNN3_QUERY_INPUT_ATTR/RKNN3_QUERY_OUTPUT_ATTR query results.

Member	Data Type	Description
index	uint32_t	Input parameter: index of the input/output tensor; the index value must be set before calling rknn3_query
name	char[RKNN3_MAX_NAME_LEN]	Name of the tensor

Member	Data Type	Description
n_dims	uint32_t	Number of dimensions
shape	uint32_t[RKNN3_MAX_DIMS]	Valid dimension array
aligned_size	uint64_t	Tensor size of the aligned shape in bytes
n_stride	uint32_t	Number of strides
stride	uint64_t[RKNN3_MAX_STRIDE_DIMS]	Strides of the tensor; e.g., strides of a 16x16 tensor are [16*16, 16, 1]
n_elems	uint32_t	Number of elements in the tensor
dtype	rknn3_tensor_type	Data type of the tensor
layout	rknn3_tensor_layout	Data layout of the tensor
qnt_type	rknn3_tensor_qnt_type	Quantization type of the tensor
qnt_info	rknn3_quant_info	Quantization information of the tensor
n_orig_dims	uint32_t	Number of original dimensions
orig_shape	uint32_t[RKNN3_MAX_DIMS]	Valid original dimension array
orig_layout	rknn3_tensor_layout	Data layout of the original tensor
core_id	int32_t	Core ID of the tensor buffer

RKNN3_TENSOR_MEM

Tensor memory information structure.

Member	Data Type	Description
virt_addr	void*	Virtual address of the tensor buffer
phys_addr	uint64_t	Physical address of the tensor buffer
fd	int32_t	File descriptor of the tensor buffer
buffer_id	uint64_t	Buffer ID of the tensor buffer, used for memory management
offset	uint64_t	Memory offset
size	uint64_t	Size of the tensor buffer
flags	uint64_t	Flags of the tensor buffer, reserved field
core_id	int32_t	NPU core ID
priv_data	void*	Private data of the tensor buffer
mem_type	rknn3_mem_type	Memory type of the tensor buffer

RKNN3_TENSOR

Structure representing an RKNN3 tensor containing memory and attribute information.

This structure contains the memory information and attributes of the tensor used in RKNN3 operations, and is the basic data structure for the RKNN3 runtime to handle tensors.

Member	Data Type	Description
mem	rknn3_tensor_mem*	Memory information of the tensor
attr	rknn3_tensor_attr*	Attributes of the tensor

5.18.2. Model and Configuration Structures

RKNN3_SDK_VERSION

SDK version information structure.

Member	Data Type	Description
api_version	char[256]	Version of the RKNN3 API
drv_version	char[256]	Version of the RKNN3 driver

RKNN3_CORE_MEM_SIZE

Structure for weight and internal memory size information for each core.

Member	Data Type	Description
core_id	int32_t	Physical NPU core ID to which the memory belongs
weight_size	uint64_t	Weight memory size in bytes
internal_size	uint64_t	Internal tensor memory size in bytes
reserved	uint8_t[32]	Reserved field

RKNN3_CUSTOM_STRING

Custom string information structure.

Member	Data Type	Description
string	char[1024]	Custom string, maximum length is 1024 bytes

RKNN3_CONFIG

Model loading control parameter structure.

Member	Data Type	Description
priority	int32_t	Run priority
run_timeout	uint32_t	Run timeout in milliseconds; 0 means no timeout
run_core_mask	uint32_t	Core mask for model execution
user_mem_weight	uint8_t	Whether weight memory is allocated by the user
user_mem_internal	uint8_t	Whether internal memory is allocated by the user
user_sram	uint8_t	Whether SRAM memory is allocated by the user
reserved	uint8_t[128]	Reserved field

RKNN3_SHAPE_INFO

RKNN3 model tensor shape information structure.

This structure contains shape information for RKNN3 model input and output tensors, including tensor attributes and shape configuration details.

Member	Data Type	Description
shape_id	int32_t	Unique ID for this shape combination
n_inputs	uint32_t	Number of input tensors
input_attrs	rknn3_tensor_attr*	Input tensor attribute array
n_outputs	uint32_t	Number of output tensors
output_attrs	rknn3_tensor_attr*	Output tensor attribute array
is_default	uint8_t	Whether this is the default shape
reserved	uint8_t[31]	Reserved field

RKNN3_SHAPE_CONFIG

Configuration structure for dynamic shape setting.

This structure contains information about shape combinations for dynamic shape inference in the RKNN3 model and the currently active shape.

Member	Data Type	Description
n_shapes	uint32_t	Number of available shape combinations
current_shape_id	int32_t	ID of the currently active shape configuration; -1 means no active shape

RKNN3_ALLOCATION_INFO

Memory allocation information structure for the RKNN3 model.

This structure provides detailed memory allocation information for different memory types (command, weight, internal, KV cache) and their distribution across NPU cores.

Member	Data Type	Description
core_id	int32_t	NPU core ID
command_mem	rknn3_tensor_mem	Command memory information
weight_mem	rknn3_tensor_mem	Weight memory information
internal_mem	rknn3_tensor_mem	Internal memory information
kvcache_mem	rknn3_tensor_mem	KV cache memory information
reserved	uint8_t[128]	Reserved field

5.18.3. Device Related Structures

RKNN3_INIT_EXTEND

Device-specific initialization information structure.

This structure is used to specify device-specific parameters during RKNN3 runtime context initialization, including device ID and reserved space for future use.

Member	Data Type	Description
device_id	char*	Input parameter indicating which device to select. Can be set to NULL if only one device is connected.
reserved	uint8_t[128]	Reserved field

RKNN3_NODE_MEM_INFO

RKNN3 device node memory information structure.

This structure provides detailed memory information for each node in an RKNN3 device, including total memory available for allocation and free memory.

Member	Data Type	Description
total	uint64_t	Total memory available for this node in bytes
free	uint64_t	Free memory available for this node in bytes

RKNN3_DEV_MEM_INFO

RKNN3 device node memory information structure.

This structure provides detailed memory information for each node in an RKNN3 device, including total memory available for allocation and free memory.

Member	Data Type	Description
node_num	uint32_t	Number of nodes in the device
sys_total	uint64_t	Total system memory of the device in bytes
sys_free	uint64_t	Free system memory of the device in bytes
node_mem_info	rknn3_node_mem_info[RKNN3_MAX_NPU_NODE_NUM]	Memory information for each node

RKNN3_DEVICE

Structure representing an RKNN3 device.

This structure contains information about a specific RKNN3 device, including its ID, type, and memory information.

Member	Data Type	Description
id	char[RKNN3_MAX_DEV_LEN]	Device ID
type	char[RKNN3_MAX_DEV_LEN]	Device type
mem_info	rknn3_dev_mem_info	Memory information of the device

RKNN3_DEVICES

Structure containing RKNN3 device information.

This structure contains the number of available RKNN3 devices.

Member	Data Type	Description
n_devices	uint32_t	Number of devices
devices	rknn3_device[RKNN3_MAX_DEVS]	Device information

5.18.4. LLM Related Structures

RKNN3_LLM_CONFIG

LLM configuration structure.

This structure contains the basic configuration required to initialize and run an LLM.

Member	Data Type	Description
chat_template	char*	Chat template string
vocab_size	uint32_t	Vocabulary size
embedding_dim	uint32_t	Embedding dimension (typically equals hidden size)
max_ctx_len	uint32_t	Maximum context length
max_position_embeddings	uint32_t	Maximum position embedding length
kvcache_store_method	rknn3_kvcache_store_method	KV cache storage method
kvcache_dtype	rknn3_kvcache_dtype	KV cache data type
kvcache_group_size	uint32_t	KV cache group quantization group size
kvcache_residual_depth	uint32_t	KV cache residual depth
model_type	char*	Model type string
task_type	rknn3_llm_task_type	LLM task type
rope_cache_host_storage	uint8_t	Whether RoPE cache is stored on the host side
n_attention_kvcache_lens	uint32_t	Number of valid attention KV cache length configurations
attention_kvcache_lens	rknn3_attention_kvcache_lens [RKNN3_MAX_ATTENTION_TYPE_NUM]	KV cache length configuration for each attention type
reserved	uint8_t[127]	Reserved field

RKNN3_VOCAB_INFO

RKNN3 model vocabulary information structure.

This structure contains vocabulary information used in the RKNN3 model, including size and special token IDs.

Member	Data Type	Description
vocab_size	int	Vocabulary size
special_bos_id	int[RKNN3_MAX_SPECIAL_BOS_ID_NUM]	Special beginning-of-sequence (BOS) token IDs
special_eos_id	int[RKNN3_MAX_SPECIAL_EOS_ID_NUM]	Special end-of-sequence (EOS) token IDs
n_special_bos_id	int	Number of special BOS token IDs
n_special_eos_id	int	Number of special EOS token IDs
linefeed_id	int	Linefeed token ID
skip_special_token	bool	Whether to skip special tokens during generation
ignore_eos_token	bool	Whether to ignore EOS tokens during generation
reserved	uint8_t[64]	Reserved field for future expansion

RKNN3_LLM_EXTEND_PARAM

LLM extended parameter structure.

Member	Data Type	Description
reserved	uint8_t[128]	Reserved field

RKNN3_SAMPLING_PARAMS

Structure defining sampling parameters for an LLM instance.

Member	Data Type	Description
top_k	int32_t	Top-K sampling parameter for token generation
top_p	float	Top-P (nucleus) sampling parameter
temperature	float	Sampling temperature, affecting the randomness of token selection
repeat_penalty	float	Penalty for repeated tokens during generation
frequency_penalty	float	Penalty for frequent tokens during generation
presence_penalty	float	Penalty based on token presence in the input

RKNN3_LLM_PARAM

Structure defining parameters for configuring an LLM instance.

Member	Data Type	Description
logits_name	char*	Name of the output logits. Only needs to be explicitly set when the model has multiple outputs; otherwise can be NULL.
max_context_len	int32_t	Maximum number of tokens in the context

Member	Data Type	Description
sampling_param	rknn3_sampling_params	Sampling parameters for token generation
vocab_info	rknn3_vocab_info	Vocabulary information
extend_param	rknn3_llm_extend_param	Extended parameters

RKNN3_LORA

Structure defining LoRA (Low-Rank Adaptation) parameters used in model fine-tuning.

Member	Data Type	Description
lora_name	char[RKNN3_MAX_NAME_LEN]	Name of the LoRA
scale	float	Scale factor applied to the LoRA

RKNN3_LLM_MULTIMODAL_TENSOR

Structure representing multimodal input (e.g., text, image, audio, and video). Top-level members

Member	Data Type	Description
name	const char*	Name of this tensor
prompt	const char*	Text prompt input
tokens	int32_t*	Token input. Mutually exclusive with prompt input; the user can provide either. Note: 1. If both inputs are provided, prompt input is used by default for inference. 2. When choosing token input, the user needs to parse text and multimodal tags into tokens according to the rules.
n_tokens	uint64_t	Number of tokens in the token input. Needs to be set when using token input.
enable_thinking	bool	Controls whether "thinking mode" is enabled
image	struct	Image substructure
audio	struct	Audio substructure
video	struct	Video substructure

image substructure members

Member	Data Type	Description
image_embed	float16*	Image embeddings (size: n_image x n_image_tokens x embedding_dim x sizeof(float16))
n_image_tokens	uint64_t	Number of tokens per image
n_image	uint64_t	Number of images
image_start	const char*	Start tag for images in multimodal input

Member	Data Type	Description
image_end	const char*	End tag for images in multimodal input
image_content	const char*	Image content placeholder tag in multimodal input
image_width	uint64_t	Image width
image_height	uint64_t	Image height

audio substructure members

Member	Data Type	Description
audio_embed	float16*	Audio embeddings (size: n_audio x n_audio_tokens x embedding_dim x sizeof(float16))
n_audio_tokens	uint64_t	Number of tokens per audio segment
n_audio	uint64_t	Number of audio segments
audio_start	const char*	Start tag for audio in multimodal input
audio_end	const char*	End tag for audio in multimodal input
audio_content	const char*	Audio content placeholder tag in multimodal input

video substructure members

Member	Data Type	Description
video_embed	float16*	Video embeddings (size: n_video x n_frame_per_video x n_frame_tokens x embedding_dim x sizeof(float16))
n_frame_tokens	uint64_t	Number of tokens per frame
n_frame_per_video	uint64_t	Number of frames per video
n_video	uint64_t	Number of videos
video_start	const char*	Start tag for video in multimodal input
video_end	const char*	End tag for video in multimodal input
video_content	const char*	Video content placeholder tag in multimodal input
frame_width	uint64_t	Video frame width
frame_height	uint64_t	Video frame height

RKNN3_LLM_TENSOR

Structure representing a tensor for large language model operations.

This structure contains the basic components needed to handle language model embeddings, including tensor name, embedding vectors, token IDs, and token count.

Member	Data Type	Description
name	const char*	Name of this tensor
prompt	const char*	Text prompt input if input_type is RKNN3_LLM_INPUT_PROMPT
embed	float16*	Pointer to embedding vectors if input_type is RKNN3_LLM_INPUT_EMBED (size: n_tokens * hidden_size)
tokens	int32_t*	Token ID array if input_type is RKNN3_LLM_INPUT_TOKEN
n_tokens	uint64_t	Number of tokens represented in the embedding
enable_thinking	bool	Controls whether "thinking mode" is enabled

RKNN3_LLM_INPUT

Structure representing different types of LLM input via a union.

Member	Data Type	Description
role	const char*	Message role: "user" (user input), "tool" (function result)
input_type	rknn3_llm_input_type	Specifies the type of input provided (e.g., token, embed, aux)
llm_input	rknn3_llm_tensor	Embedding vectors when input_type is RKNN3_LLM_INPUT_EMBED; token array when input_type is RKNN3_LLM_INPUT_TOKEN; uses the prompt field when input_type is RKNN3_LLM_INPUT_PROMPT
multimodal_input	rknn3_llm_multimodal_tensor	Multimodal input when input_type is RKNN3_LLM_INPUT_MULTIMODAL
aux_input	rknn3_aux_tensor	AUX input if input_type is RKNN3_LLM_INPUT_AUX; currently rknn3_aux_tensor and rknn3_tensor are the same type

RKNN3_LLM_INFER_PARAM

LLM inference parameter structure, defining parameters during the inference process.

Member	Data Type	Description
keep_history	int	History retention flag (1: keep history, 0: discard history)
max_new_tokens	int32_t	Maximum number of new tokens to generate
prefill_only	bool	Whether to only perform the prefill stage. When set to true, the model only performs prefill and skips the decode stage, but still samples the prefill output tokens. Defaults to false.

Member	Data Type	Description
disable_sampling	bool	Whether to disable sampling. When set to true, the model does not perform sampling, suitable for scenarios where only prefill output is needed and prefill output tokens are not desired (e.g., embedding and reranker models). Typically used with <code>prefill_only</code> (setting <code>disable_sampling</code> alone is ineffective). Defaults to false.
reserved	uint8_t[128]	Reserved field for future expansion

RKLLMRESULT

Structure representing LLM inference results.

Member	Data Type	Description
token_ids	int*	Pointer to the tokens generated by the LLM
num_tokens	int	Number of generated tokens

RKLLMCALLBACK

The RKLLMCallback structure contains callback functions for LLM operations.

Member	Data Type	Description
result_callback	LLMResultCallback	Callback function for LLM result return
result_userdata	void*	User data for LLMResultCallback
sampling_callback	LLMSamplingCallback	Optional: only used when custom sampling is needed
sampling_userdata	void*	User data for LLMSamplingCallback
tokenizer_callback	LLMTokenizerCallback	Optional: only used when a custom tokenizer is needed
tokenizer_userdata	void*	User data for LLMTokenizerCallback
embed_callback	LLMGetEmbedCallback	Optional: only used when custom embedding retrieval is needed
embed_userdata	void*	User data for LLMGetEmbedCallback
output_callback	LLMOutputCallback	Optional: only used when output tensors are needed
output_userdata	void*	User data for LLMOutputCallback
output_tensors	rknn3_tensor*	Output tensor array pointer
n_output_tensors	uint32_t	Number of output tensors
input_callback	LLMInputCallback	Optional: only used when custom input tensors need to be provided
input_userdata	void*	User data for LLMInputCallback
input_tensors_index	int32_t*	Index array of input tensors
n_input_tensors	uint32_t	Number of input tensors

RKLLMRUNSTATE

The RKLLMRunState structure contains the state information of the LLM run.

Member	Data Type	Description
n_total_tokens	uint64_t	Total number of tokens processed so far
n_reuse_tokens	uint64_t	Number of tokens in the current input that hit the KV cache, i.e., the number of KV cache reused tokens
n_max_tokens	uint64_t	Maximum number of tokens that can be processed
n_prefill_tokens	uint64_t	Number of tokens processed in the prefill stage
n_decode_tokens	uint64_t	Number of tokens generated in the decode stage
n_input_tokens	uint64_t	Number of tokens input to the model (same as n_prefill_tokens)
n_output_tokens	uint64_t	Number of tokens output by the model (n_output_tokens = n_decode_tokens + 1, where 1 is the prefill output token)
input_tokens	const int32_t*	Pointer to the token array input to the model; the user should not free this pointer
output_tokens	const int32_t*	Pointer to the model's output token array; the user should not free this pointer
kvcache_policy	rknn3_kvcache_policy	KV cache policy
n_loras_enabled	int32_t	Number of enabled LoRAs
loras_enabled	rknn3_lora*	List of enabled LoRAs

5.18.5. KV Cache Related Structures

RKNN3_KVCACHE_POLICY_PARAM

Structure defining parameters for KV cache policy. If the model contains a system prompt, n_keep and n_keep_aligned in the recurrent parameters are ignored.

Member	Data Type	Description
recurrent.n_keep	int64_t	Number of cache entries to keep under the recurrent policy; ignored if the model contains a system prompt
recurrent.n_keep_aligned	int64_t	Number of kept entries aligned to kvcache_group_size under the recurrent policy
save_checkpoint.checkpoint_start_pos	int64_t	Starting position for checkpoint saving (e.g., 0); must be aligned to 128, automatically adjusted if not aligned (no checkpoint is saved at the starting position)
save_checkpoint.checkpoint_interval	int64_t	Token interval for saving checkpoints (e.g., every 128 tokens); must be aligned to 128, automatically adjusted if not aligned

Member	Data Type	Description
save_checkpoint.max_checkpoint_count	int64_t	Maximum number of checkpoints to save from checkpoint_start_pos; must be less than (max_context_len - checkpoint_start_pos) / checkpoint_interval, otherwise automatically adjusted to the maximum available value
save_checkpoint.checkpoint_tail_overwrite	bool	When set to true, the tail/last checkpoint slot is used as a rolling overwrite slot. (max_checkpoint_count - 1) checkpoints are saved at fixed positions, and when the token position exceeds (checkpoint_start_pos + max_checkpoint_count * checkpoint_interval), subsequent checkpoints will overwrite the tail slot. For example: max_checkpoint_count=5 means 4 fixed checkpoints + 1 rolling checkpoint at the tail
reserved	uint8_t[64]	Reserved field

RKNN3_ATTENTION_KVCACHE_LEN

KV cache buffer length configuration structure, for one attention type.

Member	Data Type	Description
attention_type	rknn3_attention_type	Attention type
n_kvcache_buffer_lens	uint32_t	Number of valid KV cache buffer lengths
kvcache_buffer_lens	int32_t[RKNN3_MAX_KVCACHE_LEN_GROUPS]	KV cache buffer length array

RKNN3_KVCACHE_LEN_GROUP_INFO

KV cache length group information structure, for RKNN3_QUERY_KVCACHE_LEN_GROUP_INFO query.

Usage: First get core_number via RKNN3_QUERY_CORE_NUMBER, then allocate `rknn3_kvcache_len_group_info info[core_number]`, and query with `sizeof(rknn3_kvcache_len_group_info) * core_number` as size.

Member	Data Type	Description
core_id	int32_t	Physical ID of the NPU core
n_groups	uint32_t	Number of KV cache length groups (0 means single-group traditional mode)
active_group_id	int32_t	Currently active group ID (default 0)
kvcache_sizes	uint64_t[RKNN3_MAX_KVCACHE_LEN_GROUPS]	kvcache_sizes[group_id] is the KV cache size for that group on this core

5.18.6. Custom Operator Structures

RKNN3_CUSTOM_OP_ATTR

Custom operator attribute descriptor structure, used for querying a single operator attribute via `get_param`.

Member	Data Type	Description
name	char[RKNN3_MAX_NAME_LEN]	Operator attribute name
dtype	rknn3_tensor_type	Data type of the operator attribute
n_elems	uint32_t	Number of array elements
data	void*	Operator attribute array pointer (non-owning, only valid within the current callback scope; the caller should not cache or free it)

RKNN3_CUSTOM_OP_CONTEXT

Custom operator context structure, managed by the framework; users can store custom information via `user_data`.

Member	Data Type	Description
rknn_ctx	rknn3_context	RKNN3 context handle
priv_data	void*	Framework private data
user_data	void*	User-defined data
get_param	<code>int (*) (rknn3_custom_op_context*, const char*, rknn3_custom_op_attr*)</code>	Function pointer for querying operator parameters by attribute name; set by the framework before calling init/compute callbacks, users should not overwrite it. The returned <code>op_attr.data</code> after calling is non-owning, only valid within the current callback scope

RKNN3_CUSTOM_OP

Custom operator structure, containing operator metadata and callback function set.

Member	Data Type	Description
op_type	const char*	Custom operator type name
plugin_type	rknn3_op_plugin_type	Plugin type
target	rknn3_op_target_type	Backend execution target
version	uint32_t	Operator version number
author	const char*	Author information
description	const char*	Operator description
init	<code>int (*)(rknn3_custom_op_context*)</code>	Initialization callback (optional)
prepare	<code>int (*)(rknn3_custom_op_context*, rknn3_tensor*, uint32_t, rknn3_tensor*, uint32_t)</code>	Prepare callback (optional)

Member	Data Type	Description
compute	<code>int (*)(rknn3_custom_op_context*, rknn3_tensor*, uint32_t, rknn3_tensor*, uint32_t)</code>	Compute callback (required)
deinit	<code>int (*)(rknn3_custom_op_context*)</code>	De-initialization callback (optional)
get_output_num	<code>int (*)(rknn3_custom_op_context*)</code>	Get output count callback (post-processing plugin, optional)
get_attrs	<code>int (*)(rknn3_custom_op_context*, rknn3_tensor_attr*, uint32_t, rknn3_tensor_attr*, uint32_t)</code>	Get input/output tensor attributes callback (post-processing plugin, optional)

5.18.7. Image Processing Structures

RKNN3_IM_RECT

Image rectangle structure, used to describe a region of interest (ROI) in image processing.

Member	Data Type	Description
x	int	Top-left X coordinate
y	int	Top-left Y coordinate
width	int	Rectangle width in pixels
height	int	Rectangle height in pixels

RKNN3_IM_METADATA

Metadata structure associated with RKNN3 image memory, used for recording extra information between host and device.

Member	Data Type	Description
peer_im_mem_addr	uint64_t	Address of the peer (remote) image memory object

RKNN3_IM_MEM

Image memory information structure, describing an image buffer stored in device memory, used by image interfaces such as `rknn3_im_mem_create`, `rknn3_im_cvt_color`, etc.

Member	Data Type	Description
width	int	Image width in pixels
height	int	Image height in pixels
stride	int	Image buffer stride in bytes
format	rknn3_im_fmt	Image format; see rknn3_im_fmt
sync_to_host	bool	Whether image data needs to be written back to the host; default false
data_mem	rknn3_tensor_mem*	Pointer to the underlying tensor memory information

Member	Data Type	Description
metadata	rknn3_im_metadata	Extra metadata shared between host and device

6. RKNN3 C API Change Log

v0.3.0 -> v0.4.0 Session API Main Changes (2025-11-03)

Reference demo1: `rknn3-runtime/examples/rknn3_session_test_demo`

Reference demo2: `rknn3_model_zoo/examples/Qwen2_5_VL`

1. Added multimodal `audio/video` input types, supporting multiple modalities in arbitrary order and combination.
2. When inputting a prompt, explicitly insert modality tags, e.g.: `You need to do 3 things: 1.<audio>Convert this speech to text; 2.<image>Describe the content of this image; 3.<video>Describe the content of this video .`
3. `rknn3_llm_multimodal_tensor` structure optimized: e.g., changed `input_tensor.n_image` members to `input_tensor.image.n_image` substructure representation.
4. Added `thinking mode` toggle.
5. Enable thinking mode by setting `enable_thinking = true` (default is false).
6. Added support for configuring multiple `eos tokens` to control inference termination.
7. `params.vocab_info.special_eos_id` initialization updated to list-based assignment.
8. Added `rknn3_dump_features` interface for running the model layer by layer and exporting intermediate tensors in `.npy` format. Supports auto-allocating output tensors when not provided.

v0.4.0b0 -> v0.5.0 Main Changes (2025-11-29)

1. Added post-processing related query commands and custom operator plugin mechanism.
2. Added `RKNN3_QUERY_POSTPROCESS_*` query items.
3. Added `rknn3_register_custom_ops_plugins` and related structures and enumerations.
4. LLM callback and configuration enhancements.
5. Added `rknn3_session_set_function_tools` interface.
6. Replaced "last hidden layer callback" with `LLMOutputCallback` + `LLMOutputCallbackState` .
7. `rknn3_llm_config` added `model_type` and `task_type` fields.
8. `rknn3_kvcache_policy_param` added recurrent parameters.
9. Image/YUV related interface improvements.
10. `rknn3_im_mem_create` parameters expanded, supporting size/core_id/flags.
11. Added `rknn3_im_proc_flag` enumeration for describing image processing flows.

v0.5.0 -> v1.0.0 Main Changes (2025-12-22)

1. Added performance and memory analysis interfaces.

2. Added `rknn3_profile_ops` , `rknn3_profile_mem` .

v1.0.0 -> v1.0.4 Main Changes (2026-05-09)

1. Added streaming weight loading interface.
2. Added `rknn3_load_weight_chunk` interface, supporting segment-by-segment upload of weight data to NPU devices in streaming mode.
3. `rknn3_load_model_from_path` and `rknn3_load_model_from_data` support passing NULL weight parameter to enter streaming weight loading mode.
4. `rknn3_model_init` in streaming weight loading mode only completes model configuration; weight data must be uploaded in segments via `rknn3_load_weight_chunk` .
5. Added `rknn3_create_mem_from_fd` interface, supporting creating tensor memory objects from file descriptors.
6. API header file optimization.
7. `rknn3_api.h` removed dependency on non-essential data types.

v1.0.4 -> v1.1.0 Main Changes (2026-08-03)

1. Added KV cache grouping mechanism: `rknn3_attention_type` enumeration, `rknn3_attention_kvcache_lens` / `rknn3_kvcache_len_group_info` structures, `RKNN3_QUERY_KVCACHE_LEN_GROUP_INFO` query command, and attention KV cache length configuration fields in `rknn3_llm_config` .
2. Added interfaces: `rknn3_set_kvcache_group` , `rknn3_set_weight_mem` , `rknn3_session_pause` , `rknn3_session_resume` , `rknn3_mem_sync_range` .
3. Completed custom operator related definitions:
4. Added `rknn3_custom_op_attr` structure documentation.
5. `rknn3_custom_op_context` structure completed `get_param` function pointer field description.
6. Completed structure fields: `RKLLMCallback` (input_callback, etc.), `RKLLMRunState` (n_input_tokens, etc.), `rknn3_kvcache_policy_param` (checkpoint_tail_overwrite), `rknn3_llm_infer_param` (prefill_only/disable_sampling); `rknn3_sdk_version` field length corrected from 64 to 256; `RKNN3_MAX_LORA_NUM` corrected from 128 to 32.
7. Completed `RKNN3_QUERY_CUSTOM_STRING` query command enum value.
8. Fixed `rknn3_session_save_kvcache` interface parameter type: `char*` corrected to `const char*` .
9. `rknn3_mem_sync` interface added `MEM_SYNC_CHUNK_SIZE` environment variable description (FROM_DEVICE mode block transfer, default 2MiB).
10. Synchronized with the `rknn3_api.h` header file, correcting the following interface and data structure definition inconsistencies:
11. `rknn3_tensor_attr` structure: added the missing `n_orig_dims` , `orig_shape[RKNN3_MAX_DIMS]` , and `orig_layout` fields.
12. `rknn3_im_mem_create` interface: parameter types `width` / `height` / `size` / `core_id` corrected from `int` to `int32_t` .

13. Array-type fields in multiple structures: added missing array dimensions. The `reserved` field of

```
rknn3_core_mem_size / rknn3_config / rknn3_shape_info / rknn3_allocation_info / rknn3_init_extend / rknn3_llm_config / rknn3_vocab_info / rknn3_llm_extend_param / rknn3_llm_infer_param corrected to uint8_t[N]; rknn3_dev_mem_info's node_mem_info corrected to rknn3_node_mem_info[RKNN3_MAX_NPU_NODE_NUM]; rknn3_devices's devices corrected to rknn3_device[RKNN3_MAX_DEVS].
```